

Name-Anjali Savalkar

Roll no- 140

Usn-01fe22bcs196

1 . E-Commerce Cloud Deployment -Service model PAAS (platform as a service)

1. Project Overview

Our project focuses on building a cloud-based **e-commerce application** that provides a seamless online shopping experience. The platform enables users to register, browse products, manage their shopping cart, and for admins to manage product listings.

2. Application Functionalities

The core features implemented in the application are:

- User Authentication**
Secure registration and login system for customers.
- Product Listing and Browsing**
Products displayed with details like name, price, image, and description.
- Add/Remove to Cart**
Users can manage cart contents by adding or removing products.
- Cart Value Calculation**
Automatically calculates the total price of all items in the cart.
- Admin Panel**
Admin users can add new products to the system via a secure interface.

3. Cloud Service Provider

We chose **Render** as our cloud provider for the following reasons:

- **Ease of Use**
Simplified deployment for full-stack applications (frontend and backend).
- **Free Tier**
Offers free hosting for student projects, including web services and databases.
- **Integrated Services Used**
 - **Web Services:** Hosting Node.js + Express backend
 - **Static Sites:** Hosting the frontend (React, HTML, CSS, JS)
 - **PostgreSQL Database:** Storing user data, product info, and cart contents

4. Cloud Architecture

i) Deployment Model – Public Cloud

Our application is hosted on Render's **public cloud**, making it accessible globally. All resources are managed and maintained by Render.

- **Advantages:**
 - Cost-effective for small projects
 - Scalable and flexible
 - No infrastructure management required

ii) Service Model – Platform as a Service (PaaS)

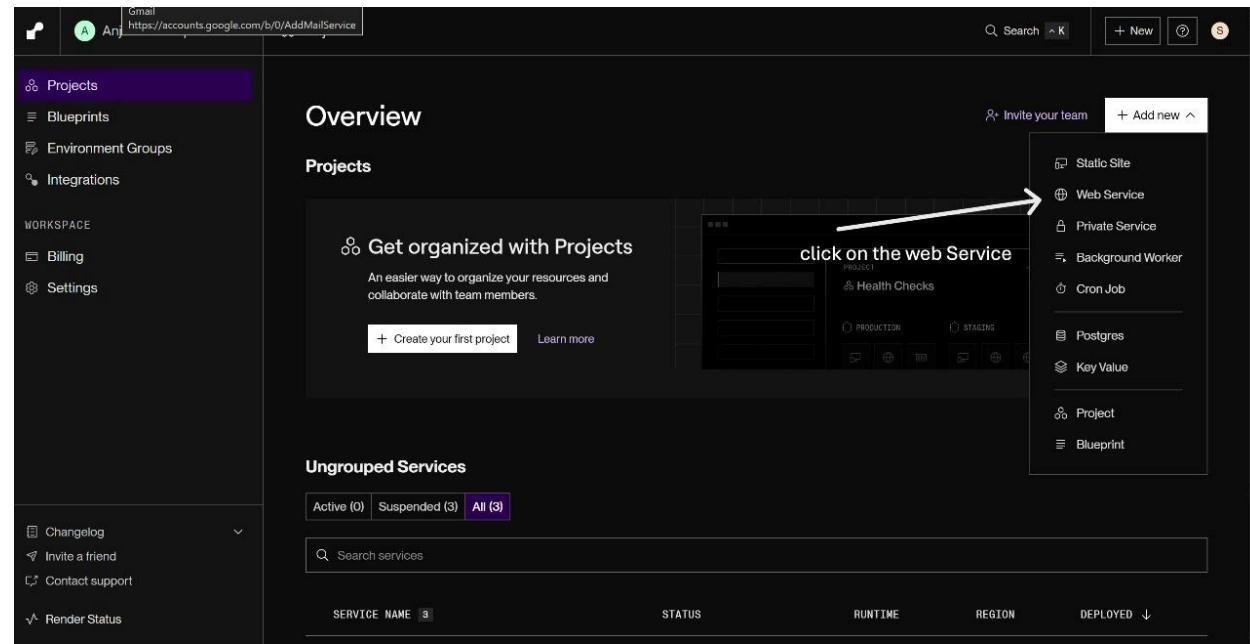
We used Render's **PaaS** capabilities, which allowed us to focus on development without managing the backend infrastructure.

- **Benefits:**
 - Easy CI/CD integration
 - Built-in support for backend/frontend hosting
 - Automated scaling and updates

5. Deployment Steps

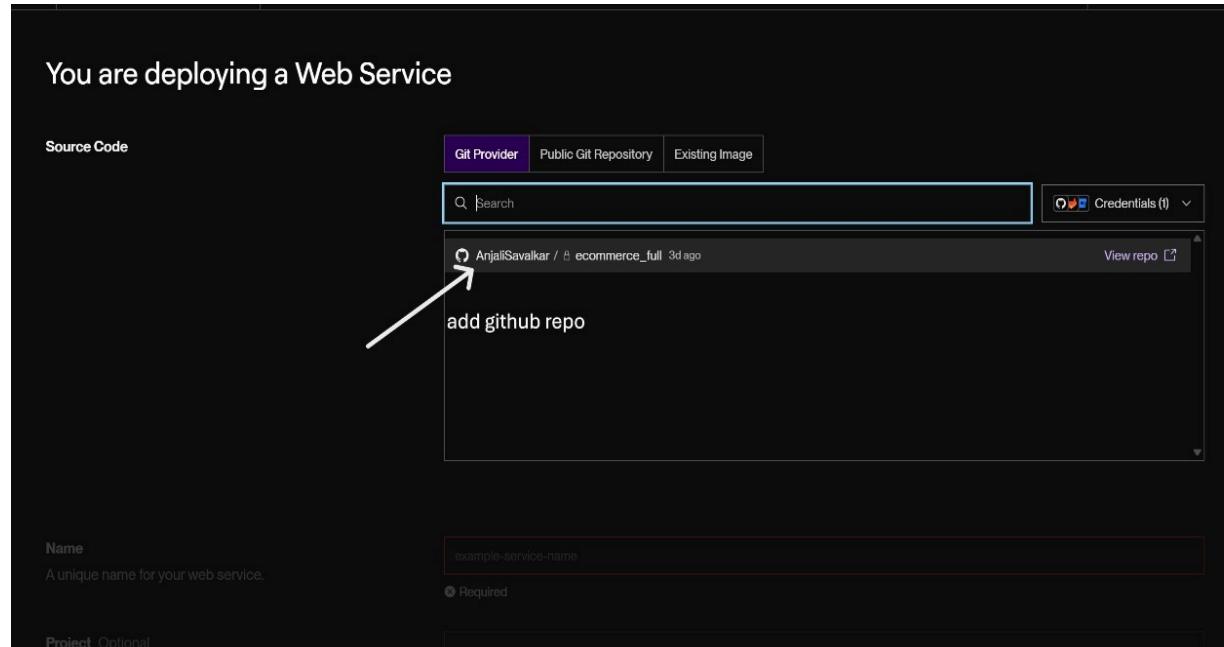
1. Render Account Setup

- Registered at [Render.com](https://render.com)
- Connected to our GitHub repository for automatic deployments

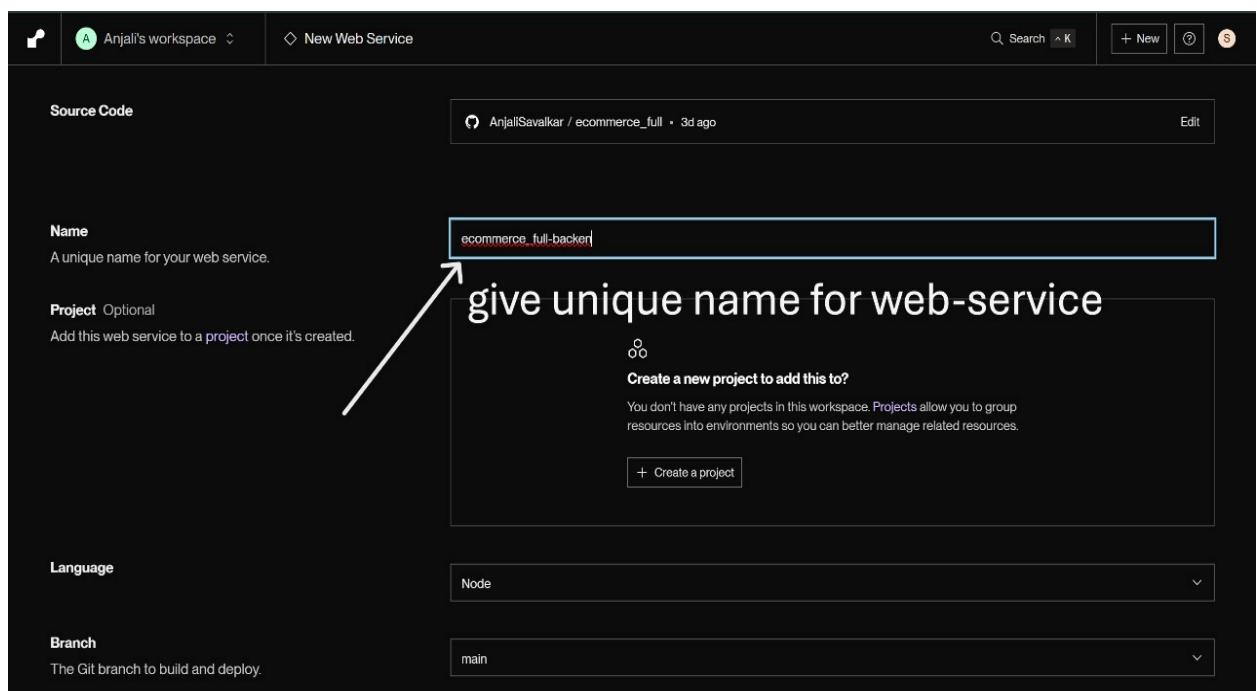


2. Frontend Deployment

- Chose "Static Site" on Render
- Linked GitHub repo (HTML/CSS/JS or React files)
- Configured build commands and publish directory
- Application hosted with a live URL

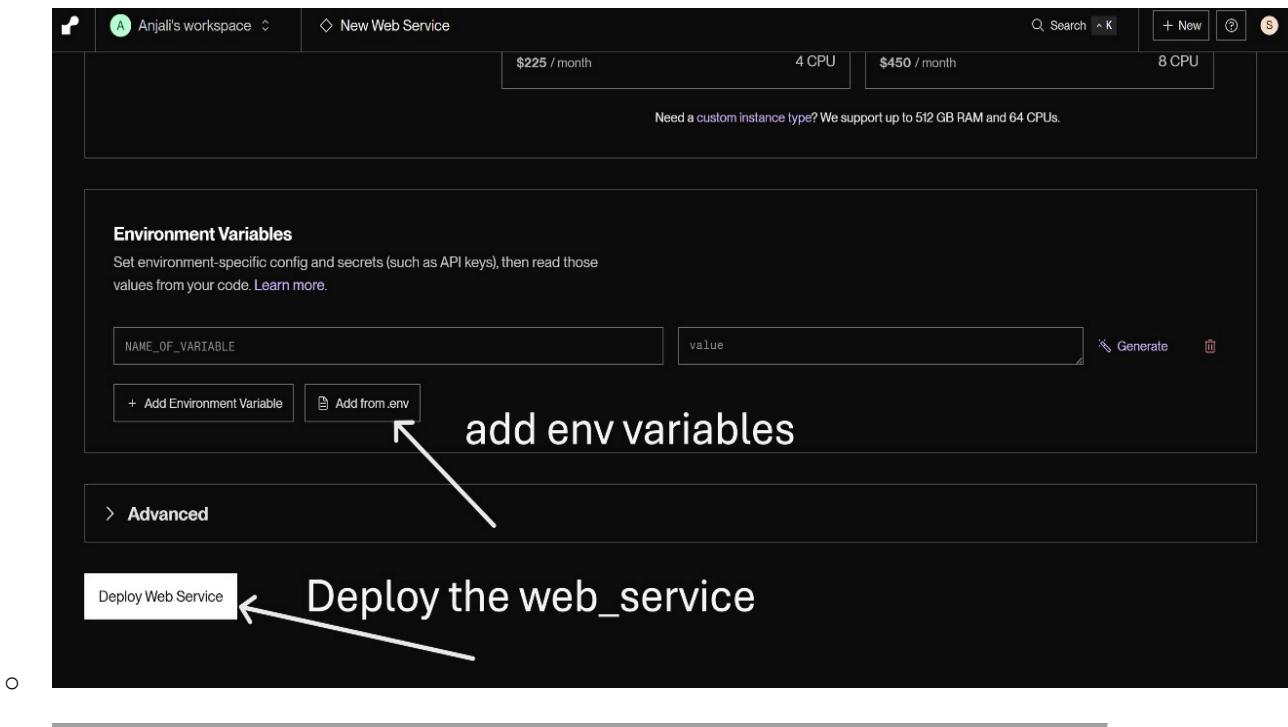


3.



4. Backend Deployment

- Set up a "Web Service" for Node.js + Express
- Configured environment variables (e.g., database credentials)
- Enabled auto-deployments from GitHub branch



Environment Variables

Set environment-specific config and secrets (such as API keys), then read those values from your code. [Learn more](#)

NAME_OF_VARIABLE value [Generate](#) [Delete](#)

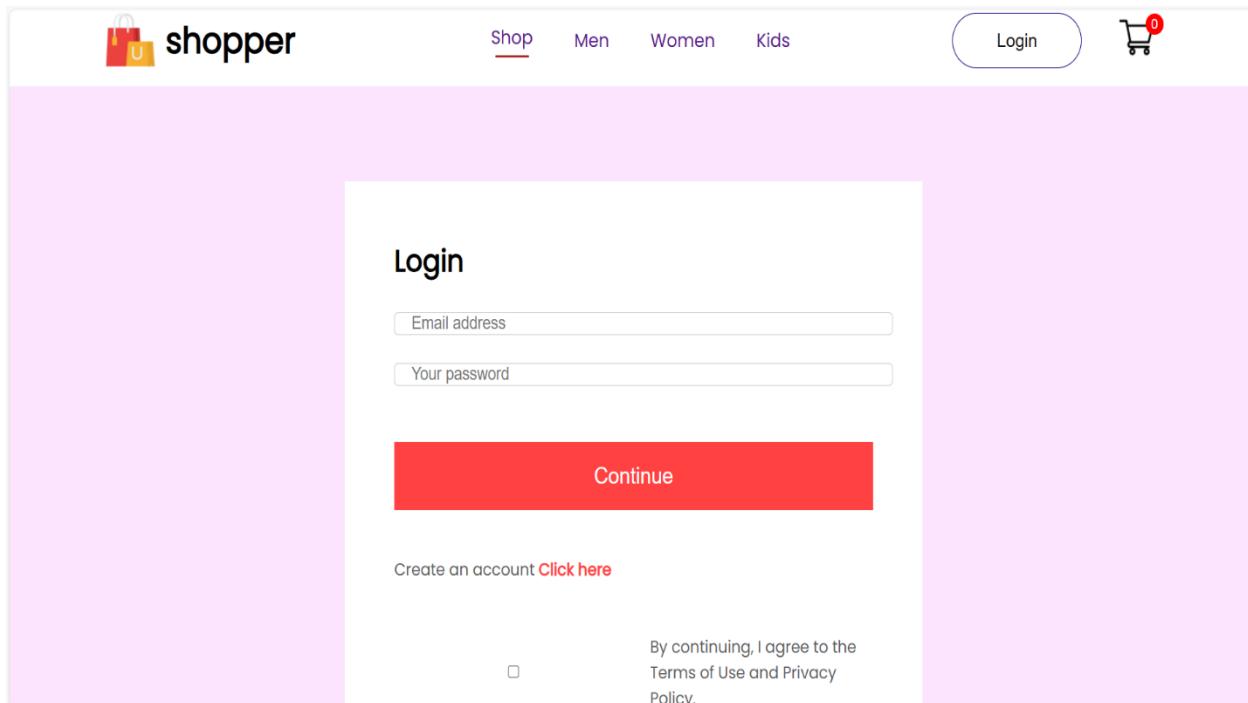
+ Add Environment Variable [Add from .env](#)

> Advanced

Deploy Web Service

6. Hosted Website

Screenshots or links of the deployed site are provided to demonstrate successful deployment and working features.



shopper

Shop Men Women Kids Login 0

Login

Email address

Your password

Continue

Create an account [Click here](#)

By continuing, I agree to the Terms of Use and Privacy Policy.

SHOPPER
Admin Panel

ecommerce-full-admin1.onrender.com says
Product added successfully!

OK

Add Product

Product title

Price

Offer Price

Product category



Add

SHOPPER
Admin Panel

All Products List

Product	Title	Old_price	new_price	Category	Remove
	Shirt	\$600	\$500	men	
	Hooded Jacket	\$799	\$699	men	
	Denim Jacket	\$800	\$850	men	
	Jacket	\$999	\$850	men	

Products Title Price Quantity Total Remove

cart totals

Subtotal		\$0
Shipping Fee		Free
Total		\$0

IF you have a promo code ,Enter it here

promo code

PROCEED TO CHECKOUT

7. Conclusion

This project showcases the development and cloud deployment of a full-stack e-commerce application using Render's PaaS platform. With features like authentication, cart management, and admin control, the application ensures a complete shopping experience while leveraging the cloud for scalability and ease of maintenance.

2 . Shopper App - Containerization of the application

Overview

The **Shopper Project** consists of three main components:

- **Frontend (shopper-frontend)**: The user interface for customers.
- **Admin Frontend (shopper-admin)**: A separate interface for managing the application.
- **Backend (shopper-backend)**: The server that handles API requests.

All components are containerized using Docker and connected via a shared Docker network.

Prerequisites

Make sure the following tools are installed:

- **Docker**
 - **Node.js**
 - **Docker Compose**
-

Dockerfiles

backend/Dockerfile

```
FROM node:18
WORKDIR /app
COPY package*.json .
RUN npm install
COPY ..
RUN npm run build
EXPOSE 5050
CMD ["npm", "start"]

frontend/Dockerfile
FROM node:18
WORKDIR /app
COPY package*.json ./
```

```
RUN npm install
COPY ..
RUN npm run build
RUN npm install -g serve
EXPOSE 5173
CMD ["serve", "-s", "dist"]
admin/vite-project/Dockerfile
FROM node:18
WORKDIR /app
COPY package*.json ./
RUN npm install
COPY ..
RUN npm run build
RUN npm install -g serve
EXPOSE 5174
CMD ["serve", "-s", "dist"]
```

Docker Compose Setup

docker-compose.yml

```
version: '3'
services:
  backend:
    build: ./backend
    image: shopper-backend:latest
    container_name: fervent_hoover
  ports:
```

```
  - "5050:5050"

networks:
  - demo

admin-frontend:
  build: ./admin/vite-project
  image: shopper-admin:latest
  container_name: reverent_raman
  ports:
    - "5174:5174"

networks:
  - demo

user-frontend:
  build: ./frontend
  image: shopper-frontend:latest
  container_name: condescending_menin
  ports:
    - "5173:5173"

networks:
  - demo

networks:
  demo:
    driver: bridge
```

Setup Instructions

1. Clone the Repository

```
git clone <repository-url>
```

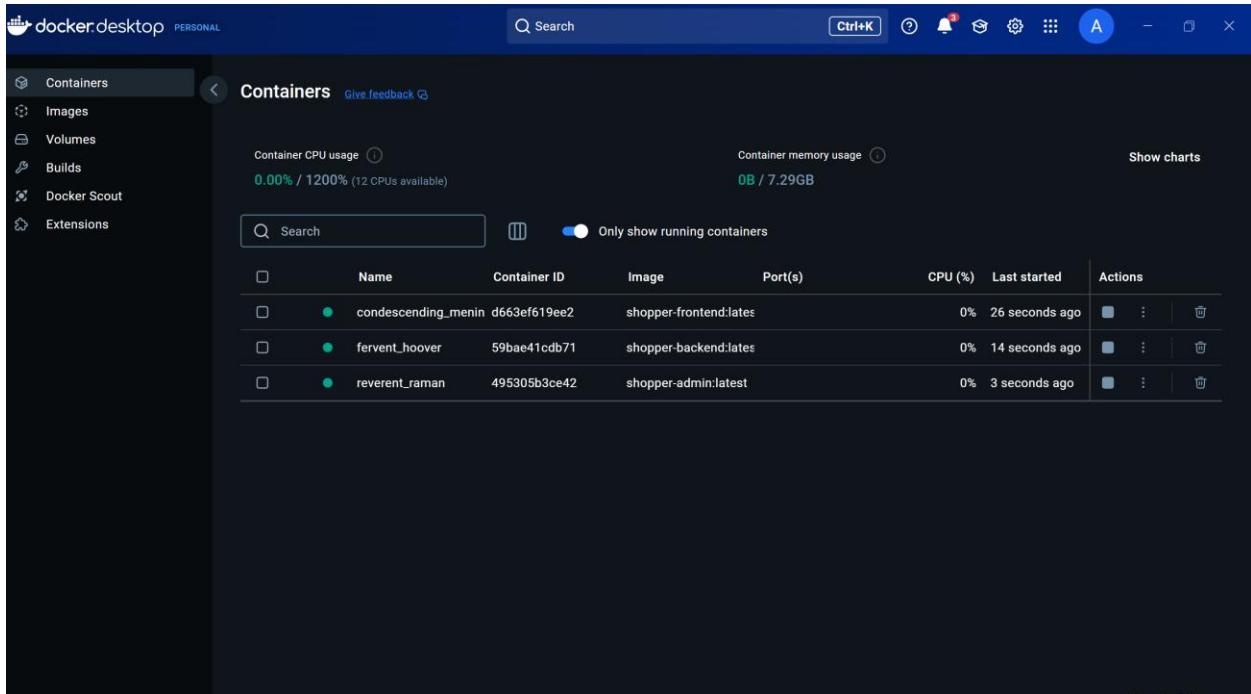
```
cd shopper
```

2. Run Docker Compose

```
docker-compose up --build
```

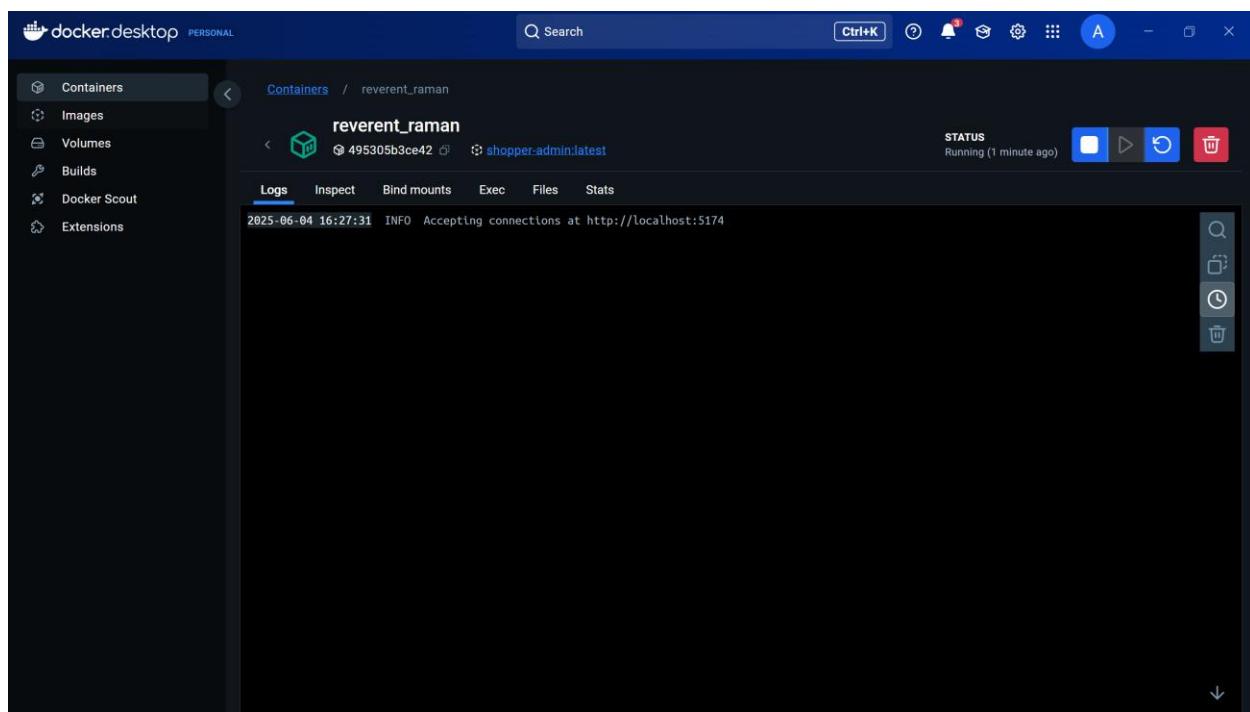
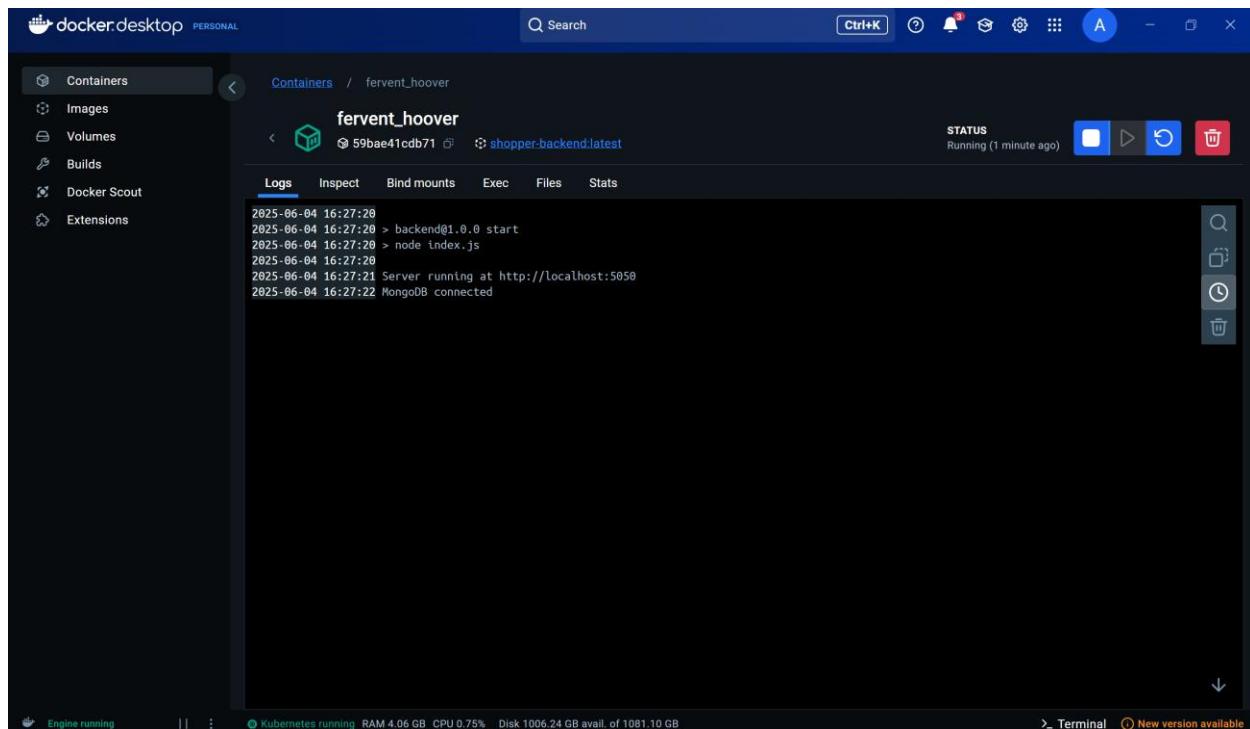
This command will build and run all three containers (backend, admin-frontend, user-frontend) and connect them to the demo network.

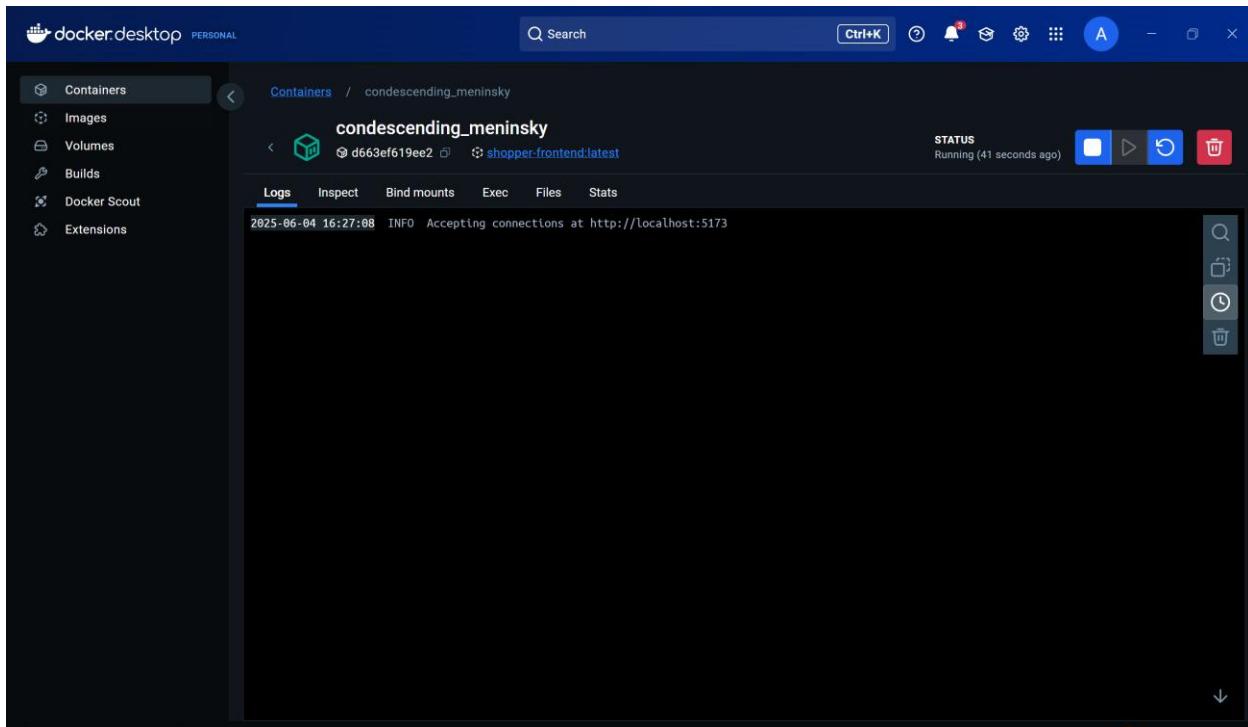
Verifying the Setup



Check Running Containers

```
docker ps
```





Access URLs

- User Frontend: <http://localhost:5173>
 - Admin Frontend: <http://localhost:5174>
 - Backend API: <http://localhost:5050>
-

3 . Docker Kubernetes – todo app

Todo App Deployment Documentation

Overview

This documentation outlines the complete process of containerizing a static frontend Todo application using Docker and deploying it on a local Kubernetes cluster via Minikube.

Tech Stack

- **Frontend:** HTML, CSS, JavaScript
 - **Web Server:** Nginx (used for serving static content)
 - **Containerization:** Docker
 - **Orchestration:** Kubernetes (Minikube)
-

Project Structure

todo-list-main/

```
├── index.html
├── styles.css
├── app.js
├── Dockerfile
└── README.md
```

[Todo App Deployment Documentation](#)

Overview

This documentation outlines the complete process of creating, containerizing, and deploying a static frontend Todo application using Docker and Kubernetes via Minikube.

Project Structure

```
todo-list-main/
├── index.html      # Main HTML page for the app
├── styles.css      # Styling for the Todo UI
├── app.js          # JavaScript to manage Todo items
├── Dockerfile       # Containerizes the app using Nginx
└── README.md       # Project overview and deployment instructions
```

Code Files

index.html

```
<!DOCTYPE html>

<html lang="en">

<head>

  <meta charset="UTF-8" />

  <meta name="viewport" content="width=device-width, initial-scale=1.0"/>

  <title>Todo App</title>

  <link rel="stylesheet" href="styles.css" />

</head>

<body>

  <div class="container">

    <h1>Todo List</h1>

    <input type="text" id="todo-input" placeholder="Add a new task" />
```

```
<button onclick="addTodo()">Add</button>

<ul id="todo-list"></ul>

</div>

<script src="app.js"></script>

</body>

</html>
```

styles.css

```
body {

  font-family: Arial, sans-serif;

  background: #f9f9f9;

  padding: 50px;

}
```

```
.container {

  background: white;

  padding: 20px;

  max-width: 500px;

  margin: auto;

  border-radius: 5px;

  box-shadow: 0 0 10px rgba(0, 0, 0, 0.1);

}
```

```
input[type="text"] {

  width: 70%;

  padding: 10px;

  margin: 10px 0;
```

```
}
```

```
button {  
  padding: 10px;  
}  
  
}
```

```
ul {  
  list-style: none;  
  padding: 0;  
}  
  
}
```

```
li {  
  padding: 8px;  
  border-bottom: 1px solid #ddd;  
}  
  
}
```

app.js

```
function addTodo() {  
  
  const input = document.getElementById('todo-input');  
  
  const text = input.value.trim();  
  
  if (text === '') return;  
  
  const list = document.getElementById('todo-list');  
  
  const li = document.createElement('li');  
  
  li.textContent = text;  
  
  list.appendChild(li);  
}  
  
}
```

```
input.value = "";  
}  


---


```

Deployment Steps

Step 1: Create the Todo App

Develop your basic HTML, CSS, and JS Todo list app.

Step 2: Dockerize the App

Create a Dockerfile:

```
FROM nginx:alpine  
COPY . /usr/share/nginx/html
```

EXPOSE 80

Build the Docker image:

```
docker build -t todo-app .
```

Run locally (optional):

```
docker run -d -p 8080:80 todo-app
```

Test: <http://localhost:8080>

Step 3: Push Image to Docker Hub (Optional)

```
docker login
```

```
docker tag todo-app your-username/todo-app
```

```
docker push your-username/todo-app
```

Step 4: Start Minikube

```
minikube start --nodes 2
```

Step 5: Create Kubernetes Manifests

deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: todo-frontend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: todo
  template:
    metadata:
      labels:
        app: todo
    spec:
      containers:
        - name: todo-container
          image: todo-app
      ports:
        - containerPort: 80
```

service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: todo-service
spec:
```

```
type: NodePort
```

```
selector:
```

```
  app: todo
```

```
ports:
```

```
  - protocol: TCP
```

```
    port: 80
```

```
    targetPort: 80
```

```
    nodePort: 30036
```

Apply the files:

```
kubectl apply -f deployment.yaml
```

```
kubectl apply -f service.yaml
```

Step 6: Access the Application

Alternative (if NodePort is not accessible):

```
kubectl port-forward service/todo-service 8080:80
```

Then open: <http://localhost:8080>

Notes

- Avoid referencing non-existent folders like ./dist in your Dockerfile unless it exists.
- You can replace **Nginx** with any other static file server if preferred.
- To view deployed pods and services:

```
kubectl get pods
```

```
kubectl get services
```

Conclusion

This guide provides a step-by-step process to containerize and deploy a simple static frontend app on Kubernetes. It's ideal for learning containerization, Docker basics, and Kubernetes deployment via Minikube.

Hpa and vpa

Horizontal Pod Autoscaling (HPA) for Todo App in Kubernetes

This guide walks you through setting up Horizontal Pod Autoscaling (HPA) for a simple Todo application deployed in a Kubernetes environment using Minikube.

Prerequisites

- Minikube installed and configured
- kubectl command-line tool installed
- Basic Todo app files: index.html, styles.css, app.js

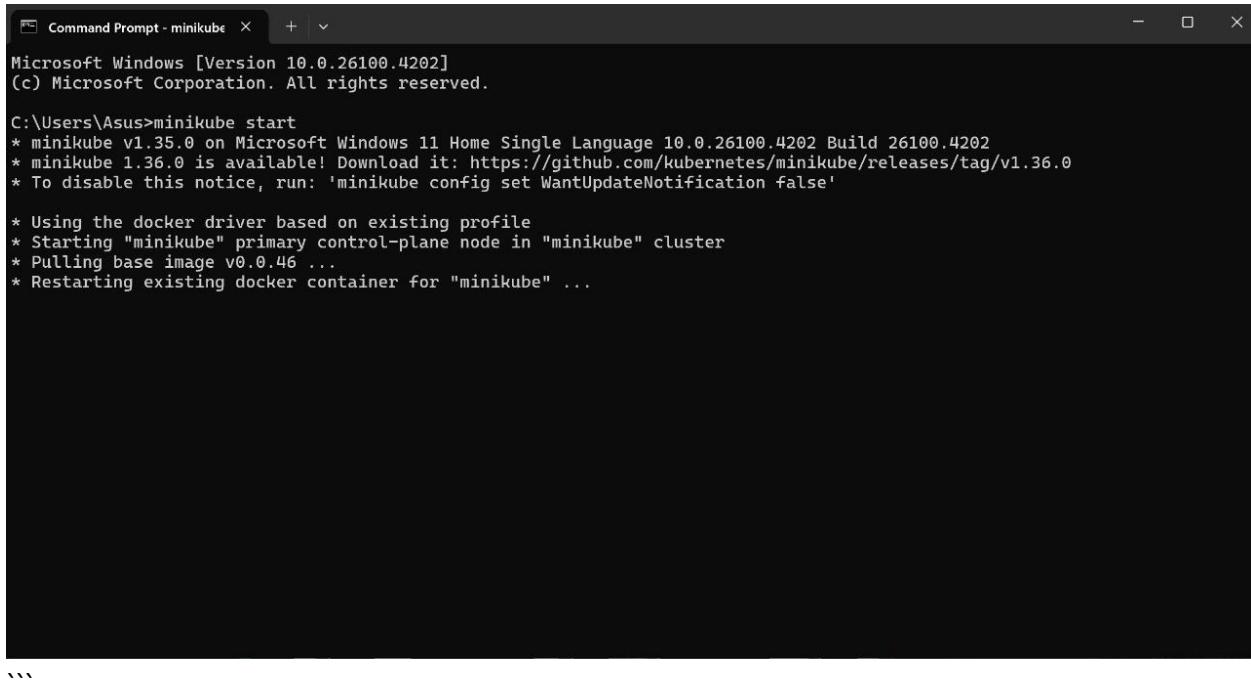
Step 1: Start Minikube with Metrics Server

HPA requires metrics to auto-scale pods. The Metrics Server provides this.

```
```bash
Start Minikube
minikube start

Enable metrics-server addon
minikube addons enable metrics-server

Verify metrics-server is running
kubectl get pods -n kube-system | grep metrics-server
```



```
Microsoft Windows [Version 10.0.26100.4202]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Asus>minikube start
* minikube v1.35.0 on Microsoft Windows 11 Home Single Language 10.0.26100.4202 Build 26100.4202
* minikube 1.36.0 is available! Download it: https://github.com/kubernetes/minikube/releases/tag/v1.36.0
* To disable this notice, run: 'minikube config set WantUpdateNotification false'

* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.46 ...
* Restarting existing docker container for "minikube" ...
```

## Step 2: Create ConfigMap for Todo App Files

Since we use a public Nginx image, create a ConfigMap to inject your app files:

```
```bash
kubectl create configmap todo-app-files --from-file=index.html --from-file=styles.css --from-file=app.js
````
```

## Step 3: Create Deployment

Create a file named todo-deployment.yaml:

```
```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: todo-app
  labels:
    app: todo-app
spec:
  replicas: 2
  selector:
    matchLabels:
```

```
app: todo-app
template:
  metadata:
    labels:
      app: todo-app
  spec:
    containers:
      - name: todo-app
        image: nginx:alpine
        imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 80
      resources:
        requests:
          cpu: 100m
          memory: 128Mi
        limits:
          cpu: 200m
          memory: 256Mi
      volumeMounts:
        - name: config-volume
          mountPath: /usr/share/nginx/html
    volumes:
      - name: config-volume
        configMap:
          name: todo-app-files
````
```

Apply the deployment:

```
```bash
kubectl apply -f todo-deployment.yaml
````
```

#### Step 4: Create Service

Create a file named todo-service.yaml to expose your app:

```
```yaml
apiVersion: v1
```

```
kind: Service
metadata:
  name: todo-app-service
spec:
  selector:
    app: todo-app
  ports:
  - port: 80
    targetPort: 80
  type: NodePort
...
```

Apply the service:

```
```bash
kubectl apply -f todo-service.yaml
````
```

Step 5: Create Horizontal Pod Autoscaler

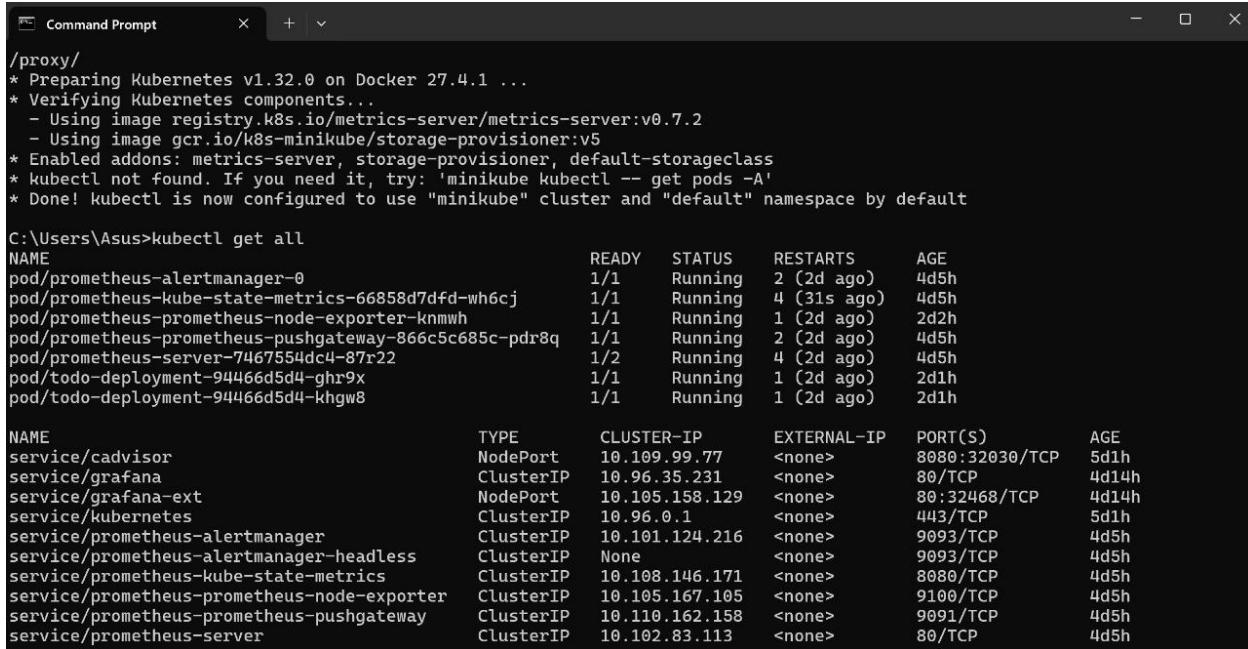
Create a file named todo-hpa.yaml to enable autoscaling based on CPU utilization:

```
```yaml
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
 name: todo-app-hpa
spec:
 scaleTargetRef:
 apiVersion: apps/v1
 kind: Deployment
 name: todo-app
 minReplicas: 2
 maxReplicas: 10
 metrics:
 - type: Resource
 resource:
 name: cpu
 target:
 type: Utilization
````
```

```
averageUtilization: 70
````
```

Apply the HPA:

```
```bash
kubectl apply -f todo-hpa.yaml
````
```



The screenshot shows a Windows Command Prompt window with the title 'Command Prompt'. The window displays two sets of command-line output. The first set is the result of the command 'kubectl get all', showing a list of pods and services. The second set is the result of the command 'kubectl get hpa', showing a list of Horizontal Pod Autoscalers. The output is as follows:

```
/proxy/
* Preparing Kubernetes v1.32.0 on Docker 27.4.1 ...
* Verifying Kubernetes components...
 - Using image registry.k8s.io/metrics-server/metrics-server:v0.7.2
 - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: metrics-server, storage-provisioner, default-storageclass
* kubectl not found. If you need it, try: 'minikube kubectl -- get pods -A'
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default

C:\Users\Asus>kubectl get all
NAME READY STATUS RESTARTS AGE
pod/prometheus-alertmanager-0 1/1 Running 2 (2d ago) 4d5h
pod/prometheus-kube-state-metrics-66858d7dfd-wh6cj 1/1 Running 4 (31s ago) 4d5h
pod/prometheus-prometheus-node-exporter-knmwh 1/1 Running 1 (2d ago) 2d2h
pod/prometheus-prometheus-pushgateway-866c5c685c-pdr8q 1/1 Running 2 (2d ago) 4d5h
pod/prometheus-server-7467554dc4-87r22 1/2 Running 4 (2d ago) 4d5h
pod/todo-deployment-94466d5d4-ghr9x 1/1 Running 1 (2d ago) 2d1h
pod/todo-deployment-94466d5d4-khgw8 1/1 Running 1 (2d ago) 2d1h

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/cadvisor NodePort 10.109.99.77 <none> 8080:32030/TCP 5d1h
service/grafana ClusterIP 10.96.35.231 <none> 80/TCP 4d14h
service/grafana-ext NodePort 10.105.158.129 <none> 80:32468/TCP 4d14h
service/kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 5d1h
service/prometheus-alertmanager ClusterIP 10.101.124.216 <none> 9093/TCP 4d5h
service/prometheus-alertmanager-headless ClusterIP None <none> 9093/TCP 4d5h
service/prometheus-kube-state-metrics ClusterIP 10.108.146.171 <none> 8080/TCP 4d5h
service/prometheus-prometheus-node-exporter ClusterIP 10.105.167.105 <none> 9100/TCP 4d5h
service/prometheus-prometheus-pushgateway ClusterIP 10.110.162.158 <none> 9091/TCP 4d5h
service/prometheus-server ClusterIP 10.102.83.113 <none> 80/TCP 4d5h

NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
todo-app-hpa Deployment/todo-app 0%/70% 2 10 2 X
```

## Step 6: Verify and Monitor HPA

Check the status of the Horizontal Pod Autoscaler:

```
```bash
kubectl get hpa todo-app-hpa
````
```

Expected output:

```
````
```

| NAME | REFERENCE | TARGETS | MINPODS | MAXPODS | REPLICAS | AGE |
|--------------|---------------------|---------|---------|---------|----------|-----|
| todo-app-hpa | Deployment/todo-app | 0%/70% | 2 | 10 | 2 | X |

```
````
```

## Step 7: Access Your Todo App

Get the URL to access the Todo app:

```
```bash
minikube service todo-app-service --url
```

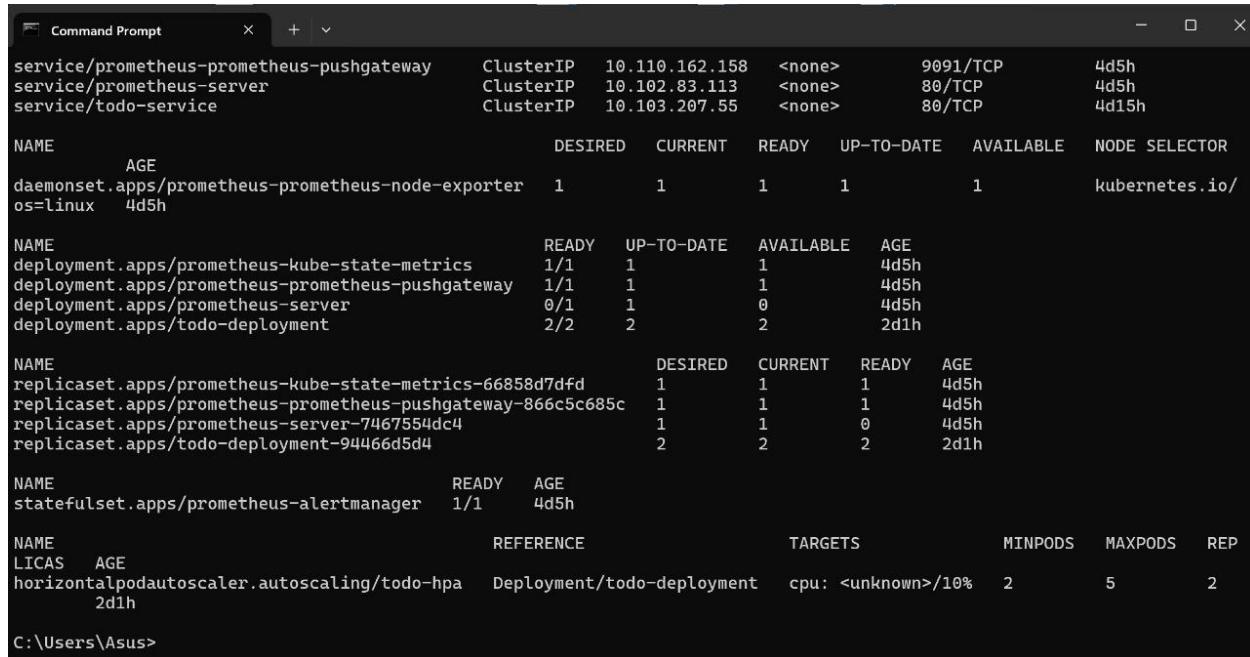
```

Or use port forwarding:

```
```bash
kubectl port-forward service/todo-app-service 8080:80
```

```

Open the URL in your browser or access via <http://localhost:8080>.



| NAME                                          | READY | UP-TO-DATE | AVAILABLE | AGE  |
|-----------------------------------------------|-------|------------|-----------|------|
| deployment.apps/prometheus-kube-state-metrics | 1/1   | 1          | 1         | 4d5h |
| deployment.apps/prometheus-pushgateway        | 1/1   | 1          | 1         | 4d5h |
| deployment.apps/prometheus-server             | 0/1   | 1          | 0         | 4d5h |
| deployment.apps/todo-deployment               | 2/2   | 2          | 2         | 2d1h |

| NAME                                                     | READY | AGE  |
|----------------------------------------------------------|-------|------|
| replicaset.apps/prometheus-kube-state-metrics-66858d7dfd | 1     | 4d5h |
| replicaset.apps/prometheus-pushgateway-866c5c685c        | 1     | 4d5h |
| replicaset.apps/prometheus-server-7467554dc4             | 1     | 4d5h |
| replicaset.apps/todo-deployment-94466d5d4                | 2     | 2d1h |

| NAME                                         | REFERENCE                  | TARGETS            | MINPODS | MAXPODS | REP |
|----------------------------------------------|----------------------------|--------------------|---------|---------|-----|
| horizontalpodautoscaler.autoscaling/todo-hpa | Deployment/todo-deployment | cpu: <unknown>/10% | 2       | 5       | 2   |

## Step 8: Generate Load to Test Autoscaling

Create `load-generator.yaml` to simulate traffic and increase CPU usage:

```
```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: load-generator
spec:
  replicas: 1
  selector:
    matchLabels:
      app: load-generator
```

```

```
template:
 metadata:
 labels:
 app: load-generator
 spec:
 containers:
 - name: load-generator
 image: busybox
 command: ["/bin/sh", "-c"]
 args:
 - while true; do
 wget -q -O http://todo-app-service;
 sleep 0.01;
 done
 ...
```

Apply the load generator:

```
```bash
kubectl apply -f load-generator.yaml
````
```

## Step 9: Monitor Autoscaling in Action

Watch the HPA scaling your pods as CPU utilization rises:

```
```bash
kubectl get hpa todo-app-hpa --watch
```

```

In another terminal, watch pod creation:

```
```bash
kubectl get pods -w
```
```
```

You should see new pods created as load increases.

Step 10: Clean Up Resources

When finished, remove all Kubernetes objects created:

```
```bash
kubectl delete -f load-generator.yaml
```

```
kubectl delete -f todo-hpa.yaml
kubectl delete -f todo-service.yaml
kubectl delete -f todo-deployment.yaml
kubectl delete configmap todo-app-files
...
```

## Troubleshooting

### Metrics Not Showing

If HPA shows <unknown>/70%:

```
```bash  
minikube addons disable metrics-server  
minikube addons enable metrics-server  
# Wait a few minutes for metrics collection to start  
...
```

Pods Not Scaling

Check current pod CPU usage:

```
```bash  
kubectl top pods
...
```

Ensure the load generator is creating enough traffic.

### Other Checks

View recent cluster events:

```
```bash  
kubectl get events --sort-by=.lastTimestamp  
...
```

Describe HPA for details:

```
```bash  
kubectl describe hpa todo-app-hpa
```

```

```
C:\Users\Asus>kubectl get hpa
NAME      REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
todo-hpa  Deployment/todo-deployment  cpu: 0%/10%  2          5          2          2d1h

C:\Users\Asus>
```

```
Command Prompt      + | -
Upper Bound:
  Cpu: 97m
  Memory: 262144k
Events: <none>

C:\Users\Asus>kubectl get hpa
NAME      REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
todo-hpa  Deployment/todo-deployment  cpu: 0%/10%  2          5          2          2d1h

C:\Users\Asus>kubectl describe hpa
Name:          todo-hpa
Namespace:     default
Labels:        <none>
Annotations:   <none>
CreationTimestamp: Mon, 02 Jun 2025 12:20:01 +0530
Reference:    Deployment/todo-deployment
Metrics:      resource cpu on pods (as a percentage of request): 0% (0) / 10%
Min replicas: 2
Max replicas: 5
Behavior:
  Scale Up:
    Stabilization Window: 0 seconds
    Select Policy: Max
  Policies:
    - Type: Pods    Value: 4    Period: 15 seconds
    - Type: Percent Value: 100   Period: 15 seconds
  Scale Down:
    Stabilization Window: 30 seconds
    - Type: Pods    Value: 2    Period: 15 seconds
```

Understanding HPA Behavior

- HPA checks metrics every 15 seconds by default.
- It averages metrics over a window to prevent rapid scaling fluctuations.
- Scale-down cooldown is longer than scale-up cooldown to avoid flapping.
- Autoscaling reactions may take several minutes to reflect.

Alternative: Using a Custom Docker Image

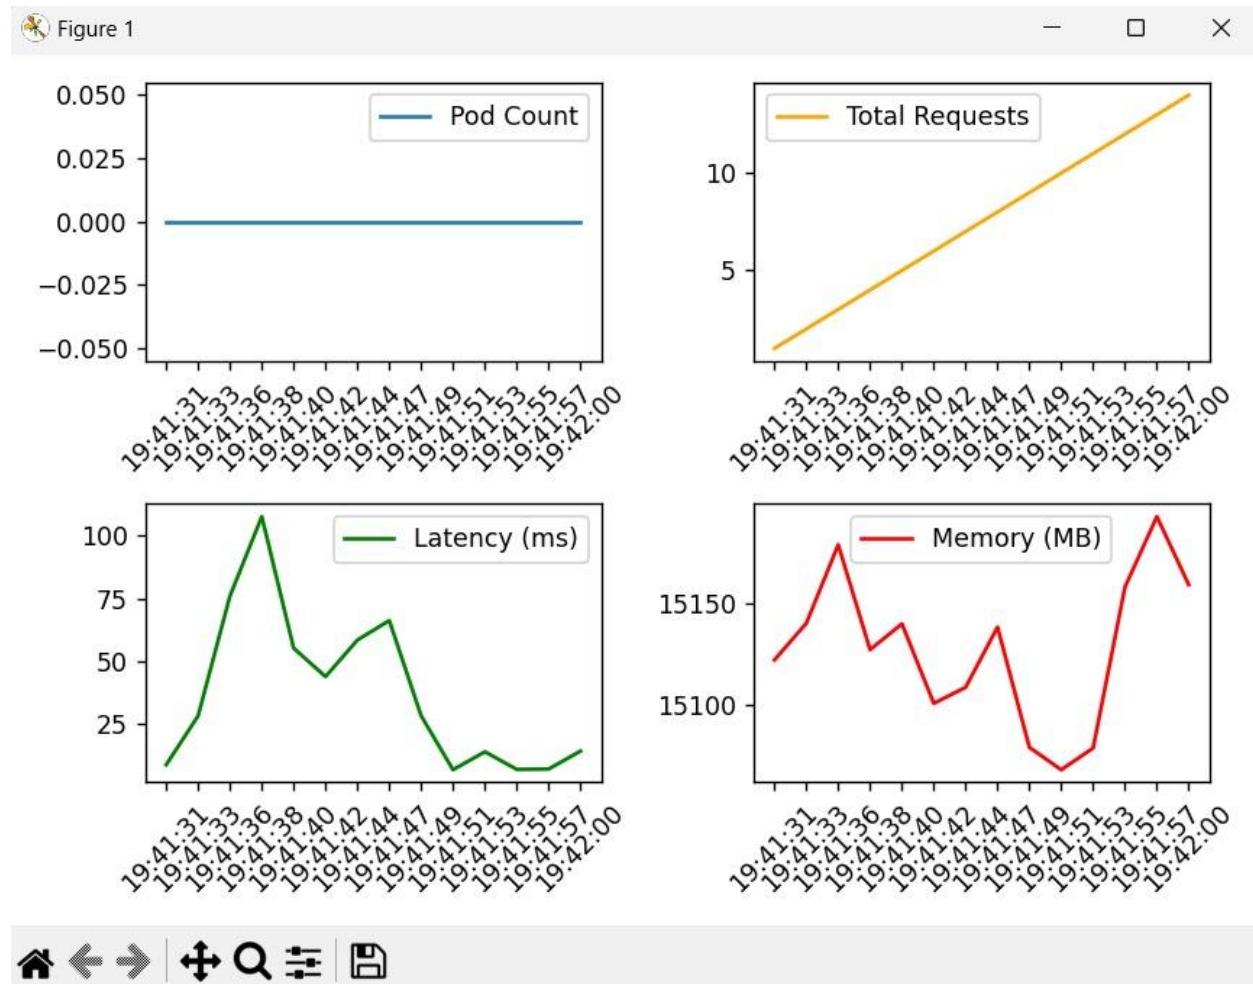
For production use, consider building your own Docker image instead of using ConfigMaps:

1. Build and push your Docker image to a container registry (e.g., Docker Hub).
2. Update your deployment to use your image and remove ConfigMap volume mounts:

```
```yaml
spec:
 containers:
 - name: todo-app
 image: yourusername/your-todo-app:latest
 # Remove the configMap volumeMounts and volumes sections
```
```

This method simplifies deployment and improves performance in production environments.

Latency and Throughput Graph in Graphana



Vpa

Vertical Pod Autoscaler (VPA) Report for TODO Application in Kubernetes

This guide explains how to generate a Vertical Pod Autoscaler (VPA) report for a TODO application running on Kubernetes. The report provides recommended CPU and memory usage for better resource optimization.

Prerequisites

- You have a TODO application running as a Kubernetes Deployment
- VPA is installed and configured on your cluster (as shown in the guide you referenced)
- Metrics server is installed and running

Step 1: Create the VPA Resource

Create a file named `vpa.yaml` for your TODO app. Here's an example VPA resource:

```
```yaml
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
 name: todo-vpa
spec:
 targetRef:
 apiVersion: "apps/v1"
 kind: Deployment
 name: todo-deployment
 updatePolicy:
 updateMode: "Auto" # Or use "Off" to only observe recommendations
````
```

Apply it:

```
```bash
kubectl apply -f vpa.yaml
````
```

Step 2: Generate Load (Optional but recommended)

To see meaningful recommendations, generate some traffic/load to the TODO application. You can use `hey` or `ab`, or a simple busybox pod like:

```
```bash
kubectl run -i --tty load-tester --image=busybox --restart=Never -- /bin/sh
````
```

```
# Inside the shell
while true; do wget -q -O- http://todo-service; done
```

```

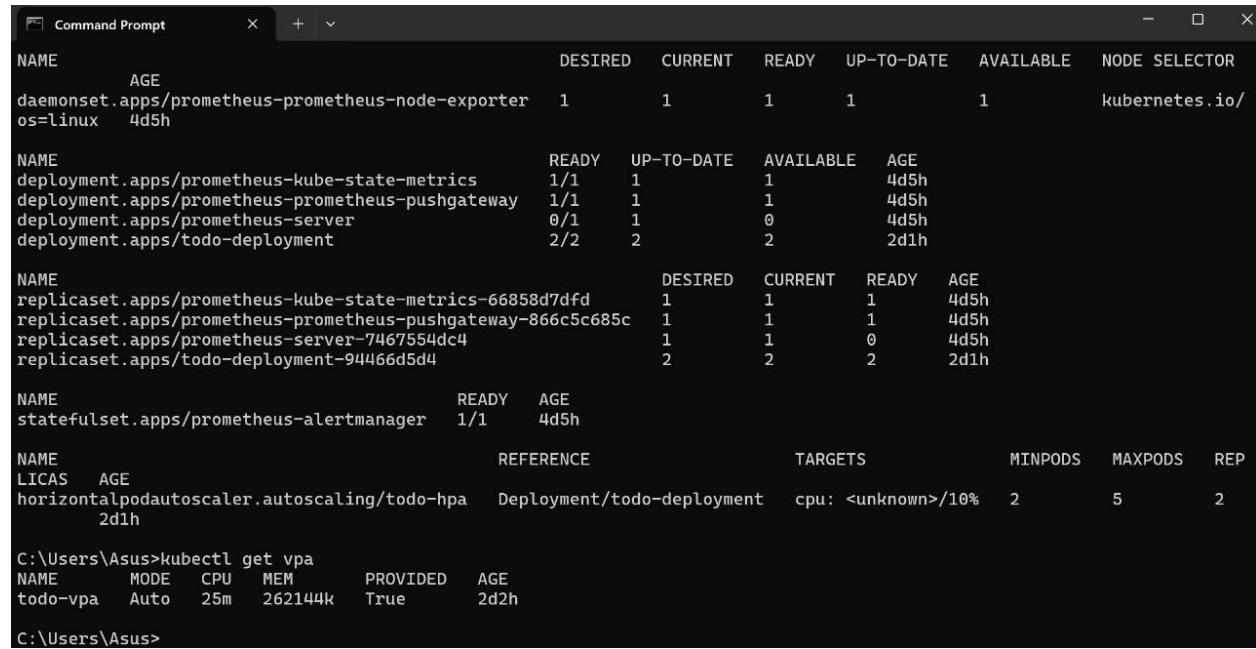
Replace `todo-service` with the correct service name.

### Step 3: View VPA Report

After some time (a few minutes of running load), run the following command to get the VPA recommendations:

```
```bash
kubectl get vpa todo-vpa -o yaml
```

```



```

Command Prompt
NAME AGE
daemonset.apps/prometheus-prometheus-node-exporter 1d5h
os=linux

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/prometheus-kube-state-metrics 1/1 1 1 4d5h
deployment.apps/prometheus-pushgateway 1/1 1 1 4d5h
deployment.apps/prometheus-server 0/1 1 0 4d5h
deployment.apps/todo-deployment 2/2 2 2 2d1h

NAME DESIRED CURRENT READY AGE
replicaset.apps/prometheus-kube-state-metrics-66858d7dfd 1 1 1 4d5h
replicaset.apps/prometheus-pushgateway-866c5c685c 1 1 1 4d5h
replicaset.apps/prometheus-server-7467554dc4 1 1 0 4d5h
replicaset.apps/todo-deployment-94466d5d4 2 2 2 2d1h

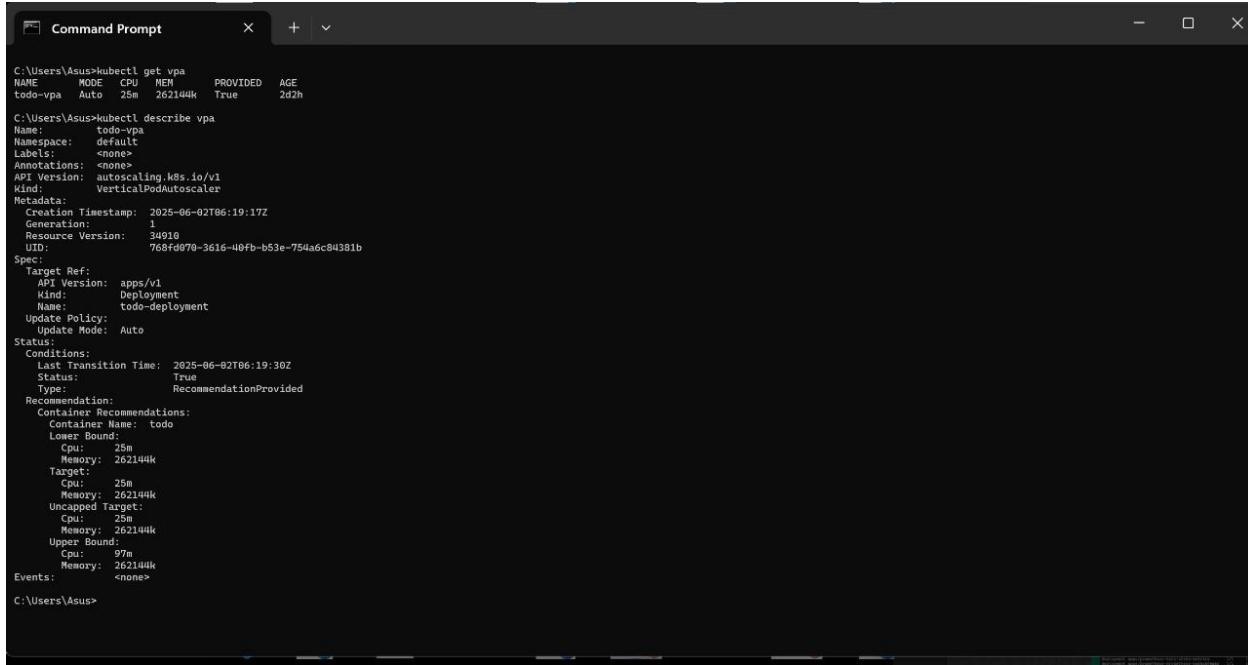
NAME READY AGE
statefulset.apps/prometheus-alertmanager 1/1 4d5h

NAME REFERENCE
LICAS AGE
horizontalpodautoscaler.autoscaling/todo-hpa Deployment/todo-deployment
 cpu: <unknown>/10% 2 5 2

C:\Users\Asus>kubectl get vpa
NAME MODE CPU MEM PROVIDED AGE
todo-vpa Auto 25m 262144k True 2d2h

C:\Users\Asus>

```



```
C:\Users\Asus>kubectl get vpa
NAME MODE CPU MEM PROVIDED AGE
todo-vpa Auto 25m 262144K True 2d2h

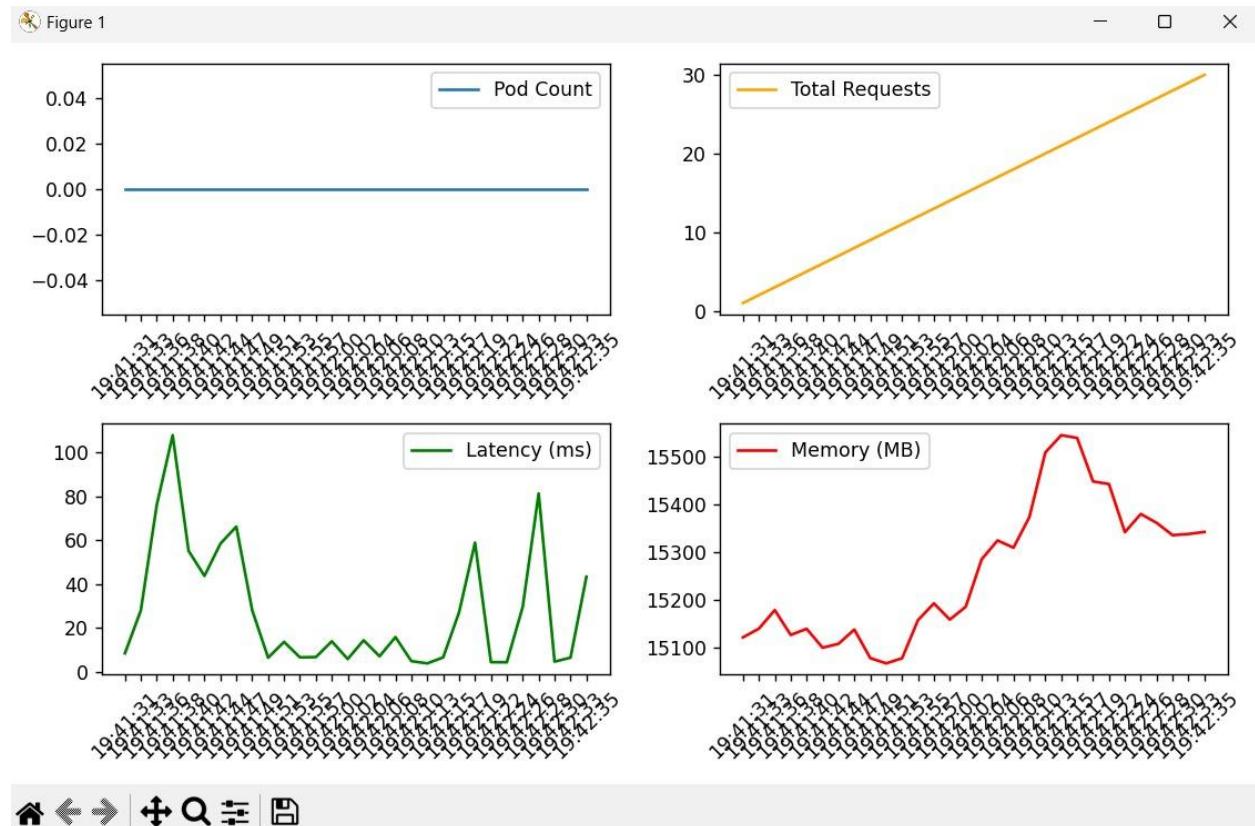
C:\Users\Asus>kubectl describe vpa
Name: todo-vpa
Namespace: default
Labels: <none>
Annotations: <none>
API Version: autoscaling.k8s.io/v1
Kind: VerticalPodAutoscaler
Metadata:
 Creation Timestamp: 2025-06-02T06:19:17Z
 Generation: 1
 Resource Version: 34910
 UID: 765fd070-3616-40fb-b53e-754a6c84381b
Spec:
 Target Ref:
 API Version: apps/v1
 Kind: Deployment
 Name: todo-deployment
 Update Policy:
 Update Mode: Auto
 Status:
 Conditions:
 Last Transition Time: 2025-06-02T06:19:30Z
 Status: True
 Type: RecommendationProvided
 Recommendation:
 Container Recommendations:
 Container Name: todo
 Lower Bound:
 Cpu: 25m
 Memory: 262144K
 Target:
 Cpu: 25m
 Memory: 262144K
 Uncapped Target:
 Cpu: 25m
 Memory: 262144K
 Upper Bound:
 Cpu: 97m
 Memory: 262144K
 Events: <none>
C:\Users\Asus>
```

## Sample Output

Here's an example snippet of what the VPA output may look like:

```
```yaml
status:
  recommendation:
    containerRecommendations:
      - containerName: todo-container
        lowerBound:
          cpu: 20m
          memory: 100Mi
        target:
          cpu: 50m
          memory: 200Mi
        upperBound:
          cpu: 100m
          memory: 400Mi
````
```

# Latency and Throughput Graph in Graphana



# **4 Microservices Cab Booking App with 4 Containers**

## ***Step 1: Define Your 4 Microservices***

***For a cab booking app, 4 common services could be:***

- 1. User Service - User login/profile UI (React or plain HTML + JS)***
- 2. Cab Search - Search available cabs UI (React or Vue)***
- 3. Booking Service - Booking confirmation UI (React or Angular)***
- 4. Payment Service - Payment UI (mock payment) (React or plain HTML + JS)***

***Since no backend, these will be UI-only apps that simulate the flow and talk to each other via mock APIs or localStorage or simple JSON.***

## ***Step 2: Setup Your Environment***

- You need Docker installed on your machine***
- You need a code editor (VSCode preferred)***
- Basic knowledge of React or HTML + JS (React for consistency)***

### **Step 3: Create Each Service as a React App**

*We'll create 4 React apps, each a simple frontend showing their UI and having buttons or forms.*

### **Step 4: Sample Code for Each Service**

#### **4.1 User Service**

*Purpose: Login screen (simulate login)*

*Create a React app user-service:*

```
npx create-react-app user-service
cd user-service
```

*Replace src/App.js with:*

```
import React, { useState } from 'react';

function App() {
 const [username, setUsername] = useState('');
 const [loggedInUser, setLoggedInUser] = useState(localStorage.getItem('user') || '');

 const login = () => {
 localStorage.setItem('user', username);
 setLoggedInUser(username);
 };

 const logout = () => {
 localStorage.removeItem('user');
 setLoggedInUser('');
 setUsername('');
 };

 return (
 <div>
 <h1>Welcome to the User Service</h1>
 <input type="text" value={username} onChange={e => setUsername(e.target.value)} />
 <button onClick={login}>Login</button>
 <button onClick={logout}>Logout</button>
 </div>
);
}

export default App;
```

```

<div style={{ padding: 20 }}>
 <h2>User Service</h2>
 {loggedInUser ? (
 <div>
 <p>Welcome, {loggedInUser}</p>
 <button onClick={logout}>Logout</button>
 </div>
) : (
 <div>
 <input
 type="text"
 placeholder="Enter username"
 value={username}
 onChange={e => setUsername(e.target.value)}
 />
 <button onClick={login} disabled={!username}>Login</button>
 </div>
)}
</div>
);
}

export default App;

```

## 4.2 Cab Search Service

*Create cab-search-service:*

```

npx create-react-app cab-search-service
cd cab-search-service

```

*Replace src/App.js with:*

```

import React, { useState } from 'react';

```

```

const cabs = [

```

```

{ id: 1, name: 'Cab A', location: 'Downtown' },
{ id: 2, name: 'Cab B', location: 'Airport' },
{ id: 3, name: 'Cab C', location: 'Suburbs' },
];

function App() {
 const [selectedCab, setSelectedCab] = useState(null);

 const bookCab = (cab) => {
 localStorage.setItem('selectedCab', JSON.stringify(cab));
 alert(`Cab ${cab.name} selected. Proceed to Booking Service.`);
 };

 return (
 <div style={{ padding: 20 }}>
 <h2>Cab Search Service</h2>

 {cabs.map(cab => (
 <li key={cab.id}>
 {cab.name} - {cab.location}' '
 <button onClick={() => bookCab(cab)}>Select</button>

)));

 </div>
);
}

export default App;

```

### 4.3 Booking Service

*Create booking-service:*

```

npx create-react-app booking-service
cd booking-service

```

Replace `src/App.js` with:

```
import React, { useState, useEffect } from 'react';

function App() {
 const [cab, setCab] = useState(null);
 const [bookingConfirmed, setBookingConfirmed] = useState(false);

 useEffect(() => {
 const selectedCab = localStorage.getItem('selectedCab');
 if (selectedCab) {
 setCab(JSON.parse(selectedCab));
 }
 }, []);

 const confirmBooking = () => {
 localStorage.setItem('bookingConfirmed', 'true');
 setBookingConfirmed(true);
 };

 return (
 <div style={{ padding: 20 }}>
 <h2>Booking Service</h2>
 {cab ? (
 <>
 <p>Selected Cab: {cab.name}</p>
 {!bookingConfirmed ? (
 <button onClick={confirmBooking}>Confirm Booking</button>
) : (
 <p>Booking Confirmed!</p>
)}
 </>
) : (
 <p>No cab selected. Please select a cab from the Cab Search Service.</p>
)}
 </div>
);
}

export default App;
```

```
};

}

export default App;
```

#### 4.4 Payment Service

*Create payment-service:*

```
npx create-react-app payment-service
cd payment-service
```

*Replace src/App.js with:*

```
import React, { useState, useEffect } from 'react';

function App() {
 const [bookingConfirmed, setBookingConfirmed] = useState(false);
 const [paymentDone, setPaymentDone] = useState(false);

 useEffect(() => {
 const booking = localStorage.getItem('bookingConfirmed');
 if (booking === 'true') {
 setBookingConfirmed(true);
 }
 }, []);

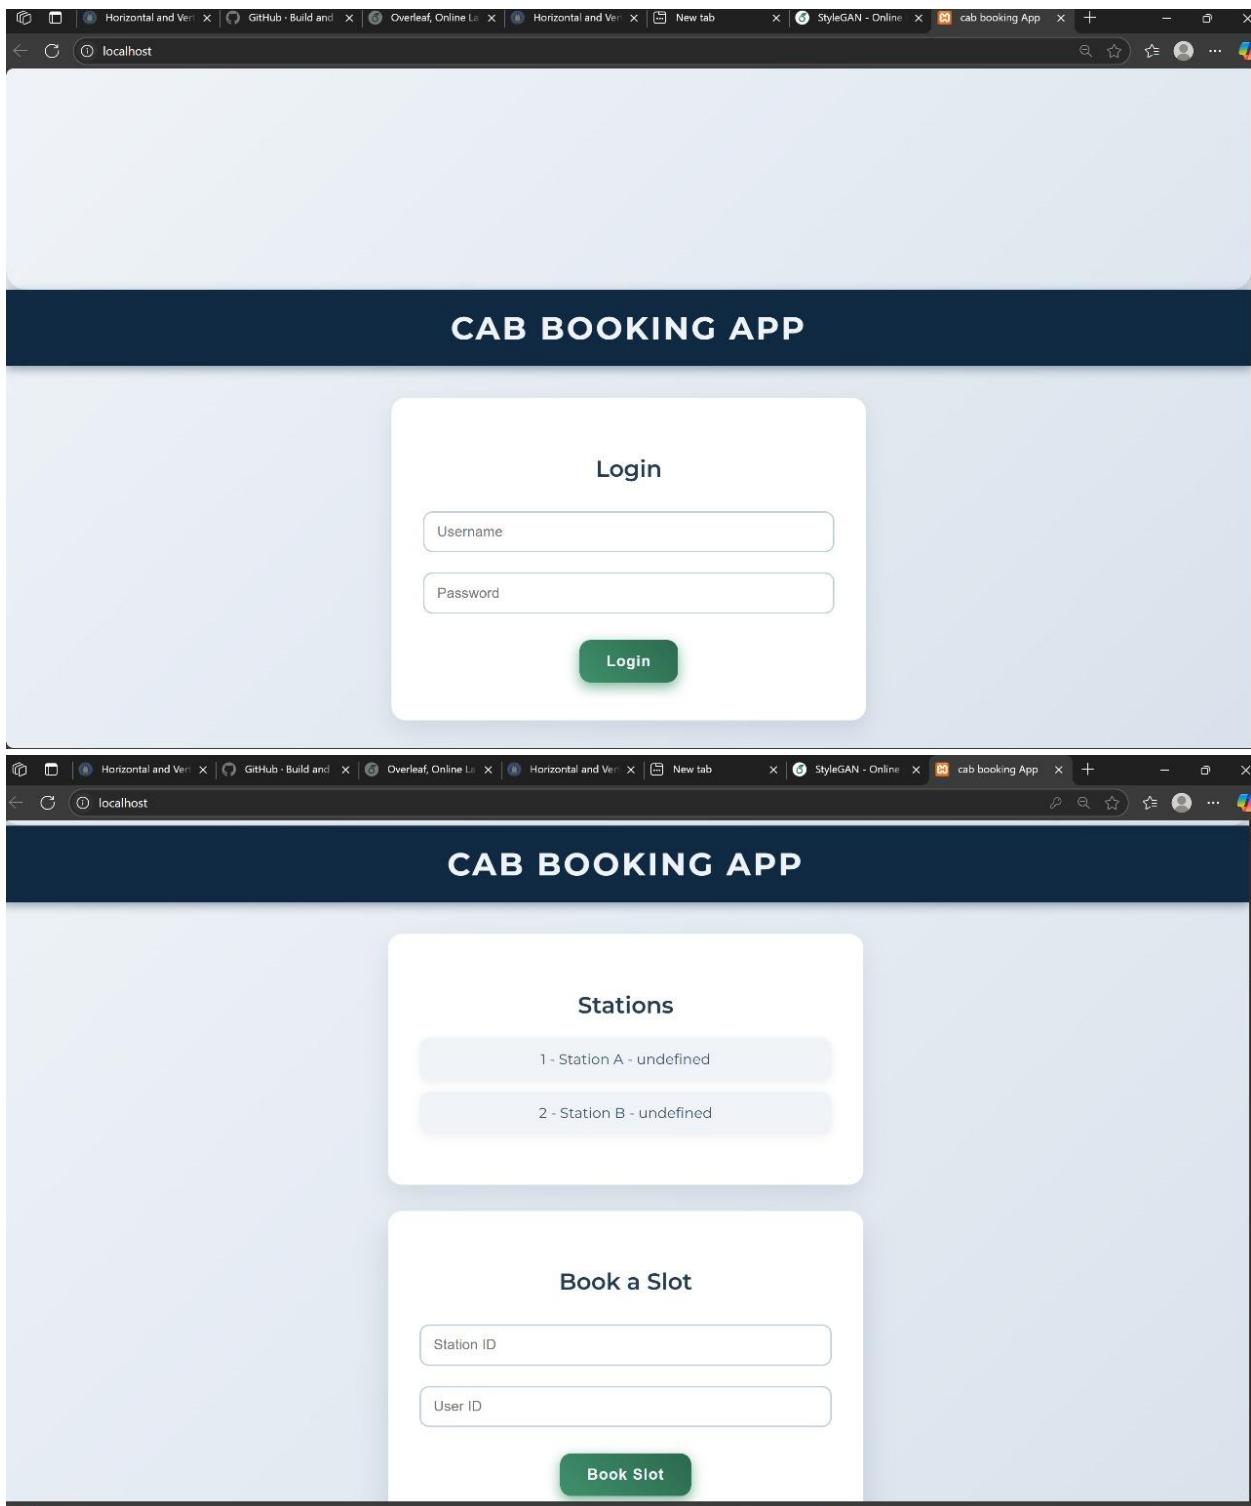
 const makePayment = () => {
 setPaymentDone(true);
 alert('Payment successful! Thank you.');
 };

 return (
 <div style={{ padding: 20 }}>
 <h2>Payment Service</h2>
 {!bookingConfirmed ? (

```

```
<p>Please confirm booking first in Booking Service.</p>
):(
<>
{!paymentDone ? (
 <button onClick={makePayment}>Pay Now</button>
):(
 <p>Payment completed. Have a safe journey!</p>
)}
</>
)}
</div>
);
}

export default App;
```



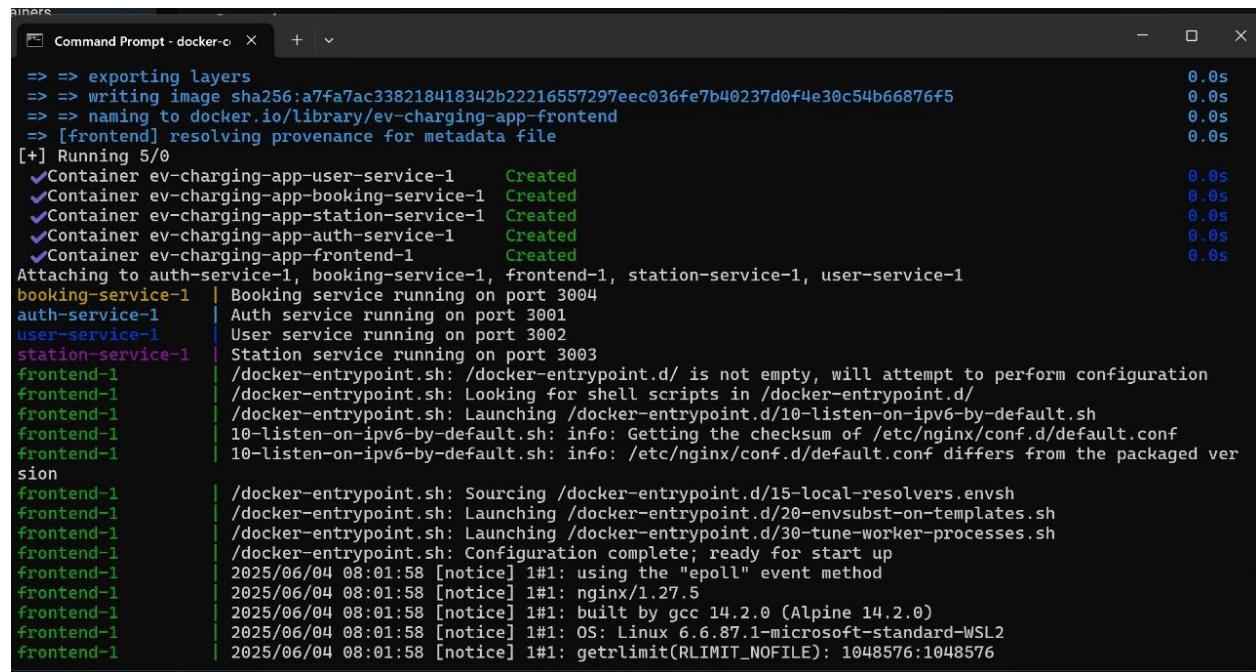
### Step 5: Dockerize Each Service

**Create a Dockerfile inside each service folder (user-service/Dockerfile, etc.)**

**Dockerfile for React apps:**

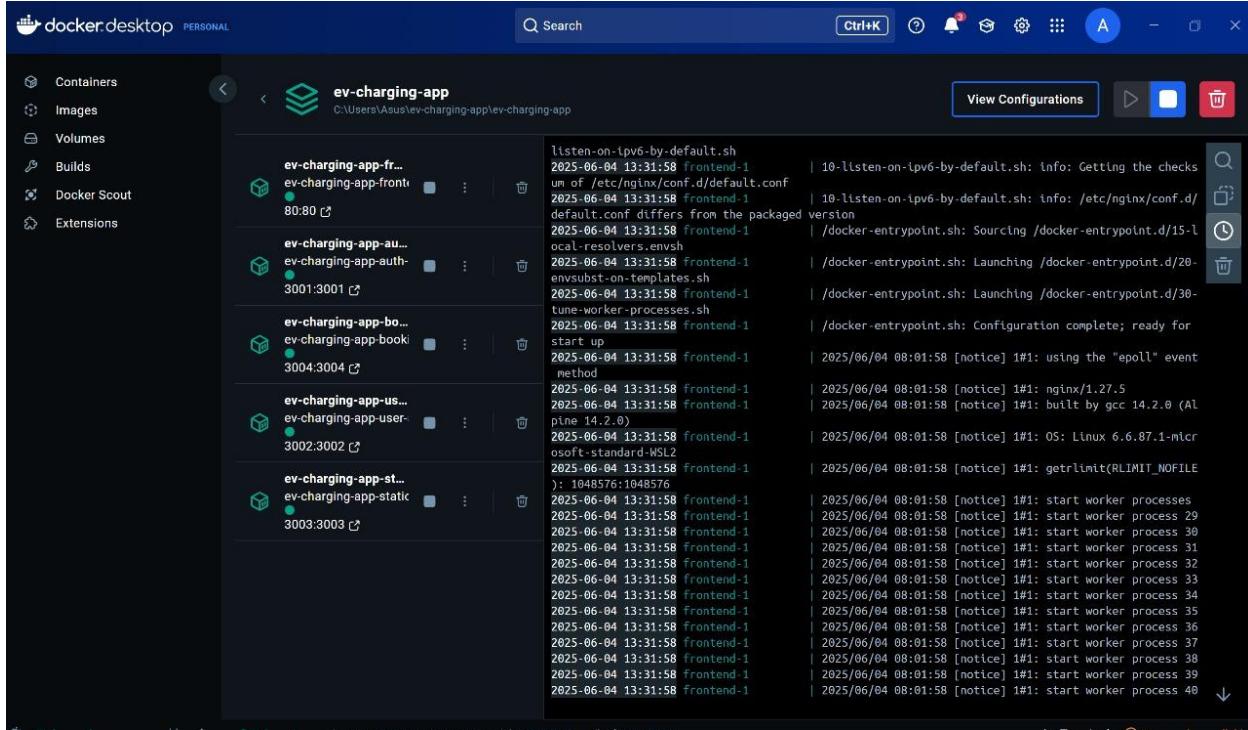
```
Dockerfile
FROM node:18-alpine
WORKDIR /app
COPY package*.json .
RUN npm install
COPY ..
RUN npm run build

FROM nginx:alpine
COPY --from=0 /app/build /usr/share/nginx/html
EXPOSE 80
CMD ["nginx", "-g", "daemon off;"]
```



The screenshot shows a Windows Command Prompt window titled "Command Prompt - docker-c" with the following log output:

```
=> => exporting layers
=> => writing image sha256:a7fa7ac338218418342b22216557297eec036fe7b40237d0f4e30c54b66876f5
=> => naming to docker.io/library/ev-charging-app-frontend
=> [frontend] resolving provenance for metadata file
[+] Running 5/0
 ✓ Container ev-charging-app-user-service-1 Created 0.0s
 ✓ Container ev-charging-app-booking-service-1 Created 0.0s
 ✓ Container ev-charging-app-station-service-1 Created 0.0s
 ✓ Container ev-charging-app-auth-service-1 Created 0.0s
 ✓ Container ev-charging-app-frontend-1 Created 0.0s
Attaching to auth-service-1, booking-service-1, frontend-1, station-service-1, user-service-1
booking-service-1 | Booking service running on port 3004
auth-service-1 | Auth service running on port 3001
user-service-1 | User service running on port 3002
station-service-1 | Station service running on port 3003
frontend-1 | /docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
frontend-1 | /docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
frontend-1 | /docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
frontend-1 | 10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
frontend-1 | 10-listen-on-ipv6-by-default.sh: info: /etc/nginx/conf.d/default.conf differs from the packaged version
frontend-1 | /docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
frontend-1 | /docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
frontend-1 | /docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
frontend-1 | /docker-entrypoint.sh: Configuration complete; ready for start up
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: using the "epoll" event method
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: nginx/1.27.5
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: built by gcc 14.2.0 (Alpine 14.2.0)
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: OS: Linux 6.6.87.1-microsoft-standard-WSL2
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
```



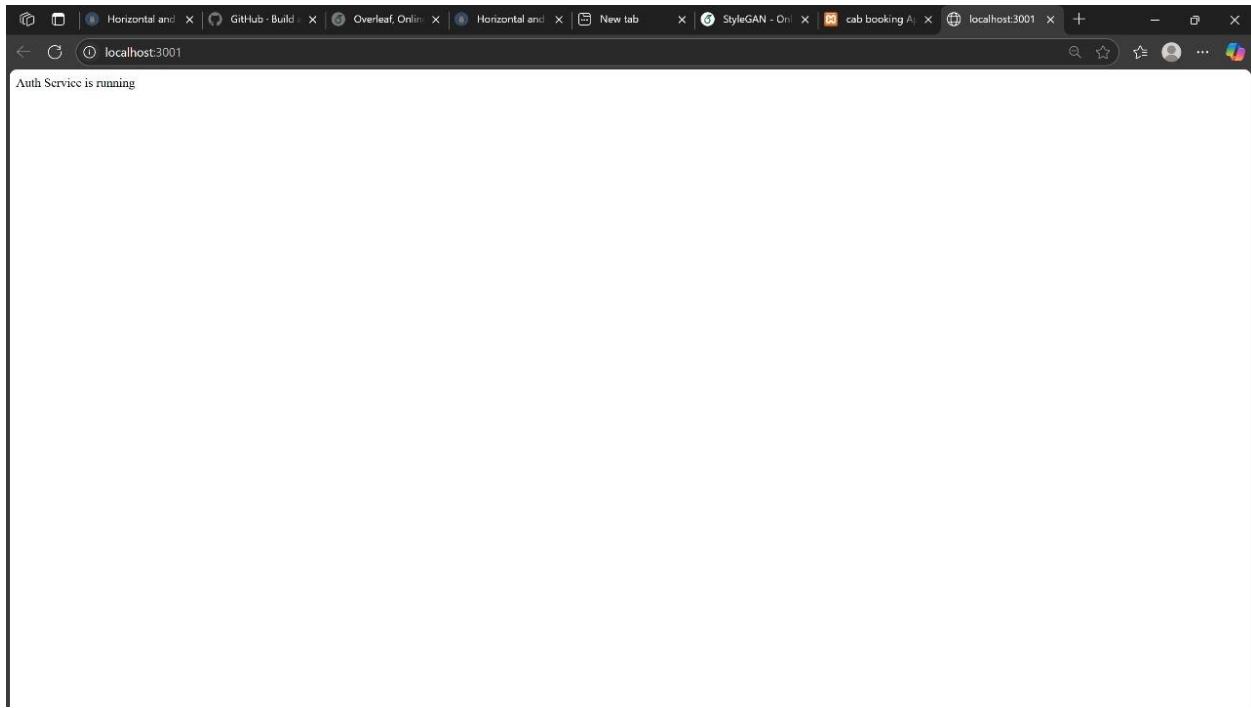
## Step 6: Build and Run Docker Containers Locally

**From inside each service folder:**  
**docker build -t <service-name> .**

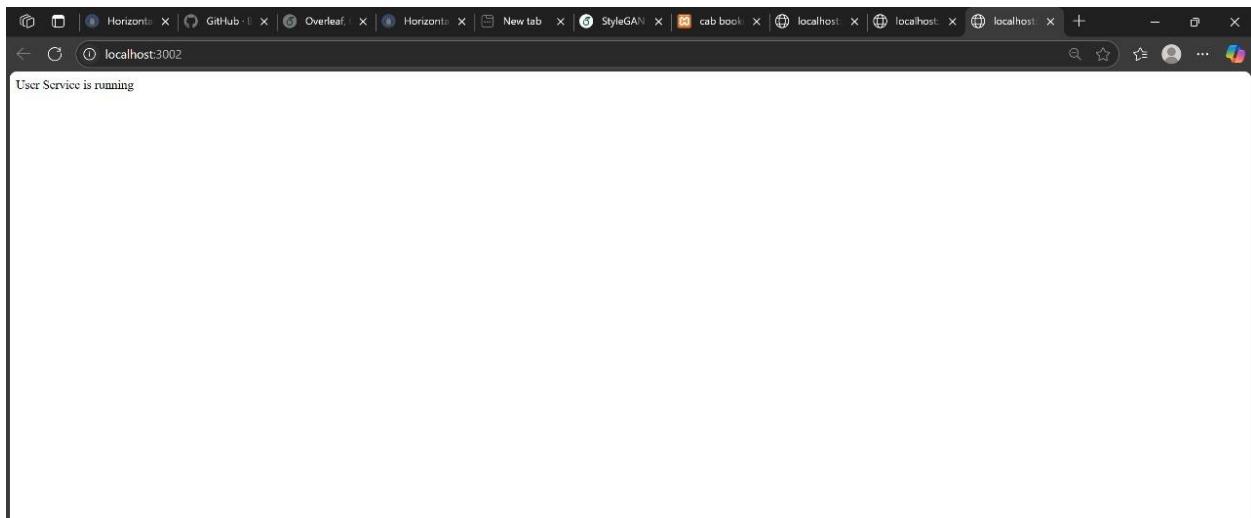
**Example:**  
**docker build -t user-service .**

**Run container:**  
**docker run -d -p <port>:80 --name <service-name> user-service**

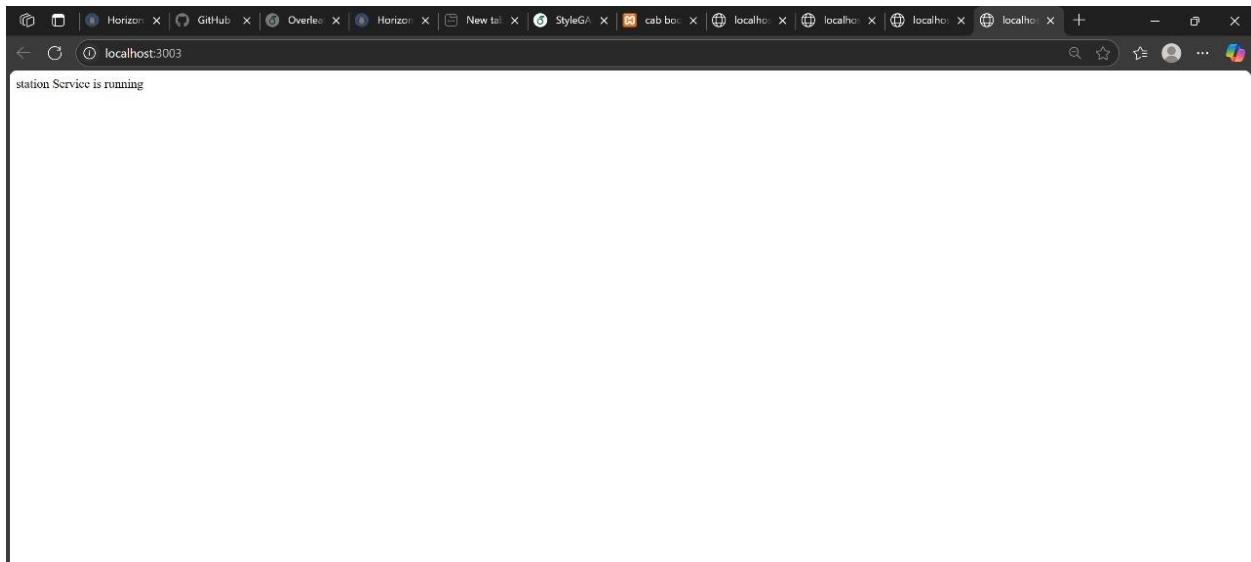
**Example:**  
**docker run -d -p 3001:80 --name user-service user-service**



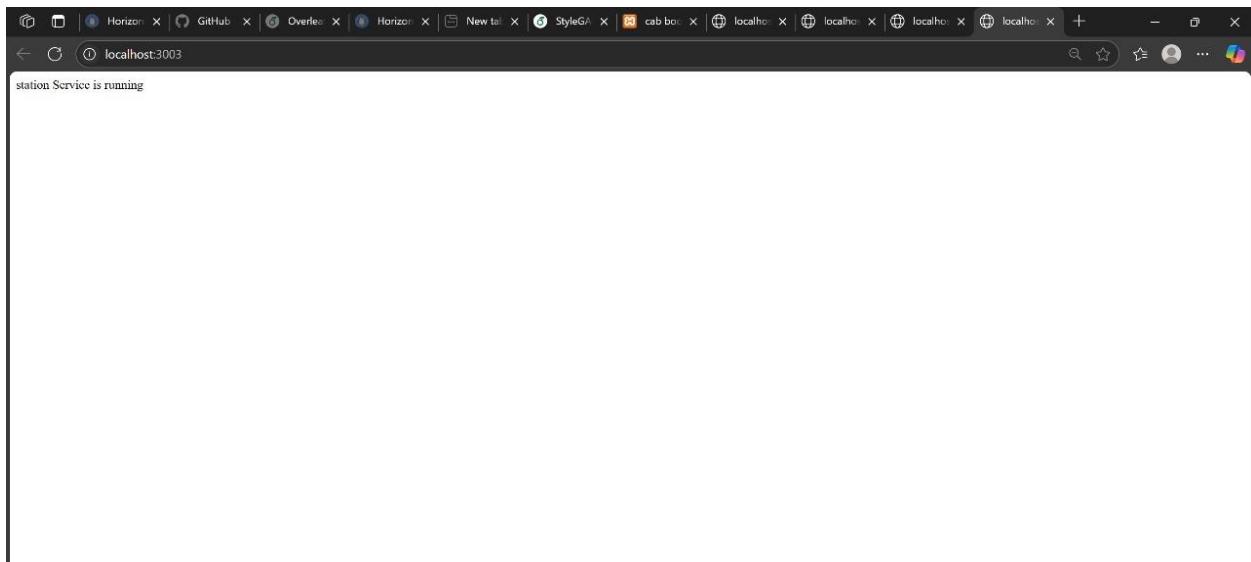
***docker run -d -p 3002:80 --name cab-search-service cab-search-service***



***docker run -d -p 3003:80 --name booking-service booking-service***



```
docker run -d -p 3004:80 --name payment-service payment-service
```



### **Step 7: Access Services**

- **User Service:** `http://localhost:3001`
- **Cab Search:** `http://localhost:3002`
- **Booking Service:** `http://localhost:3003`
- **Payment Service:** `http://localhost:3004`

### **Step 8: Test Flow**

- *Login at User Service (simulate login)*
- *Go to Cab Search (choose cab)*
- *Go to Booking (confirm booking)*
- *Go to Payment (pay)*

*Optional Step 9: Compose with Docker Compose*

*Create a docker-compose.yml in the root folder:*

```
version: '3'
services:
 user-service:
 build: ./user-service
 ports:
 - "3001:80"
 cab-search-service:
 build: ./cab-search-service
 ports:
 - "3002:80"
 booking-service:
 build: ./booking-service
 ports:
 - "3003:80"
 payment-service:
 build: ./payment-service
 ports:
 - "3004:80"
```

*Then run:*

*docker-compose up --build*

```
Command Prompt - docker-c X + v - □ ×

frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: start worker process 39
frontend-1 | 2025/06/04 08:01:58 [notice] 1#1: start worker process 40
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:23 +0000] "GET / HTTP/1.1" 200 1036 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:23 +0000] "GET /style.css HTTP/1.1" 200 2961 "http://localhost/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:23 +0000] "GET /script.js HTTP/1.1" 304 0 "http://localhost/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:23 +0000] "GET /favicon.ico HTTP/1.1" 404 555 "http://localhost/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
frontend-1 | 2025/06/04 08:03:23 [error] 29#29: *1 open() "/usr/share/nginx/html/favicon.ico" failed (2: No such file or directory), client: 172.22.0.1, server: , request: "GET /favicon.ico HTTP/1.1", host: "localhost", referrer: "http://localhost/"
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:51 +0000] "POST /api/auth/Login HTTP/1.1" 200 18 "http://localhost/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
frontend-1 | 172.22.0.1 - - [04/Jun/2025:08:03:51 +0000] "GET /api/station/stations HTTP/1.1" 200 57 "http://localhost/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/137.0.0.0 Safari/537.36 Edg/137.0.0.0" "-"
Gracefully stopping... (press Ctrl+C again to force)
[+] Stopping 1/5
 ✓ Container ev-charging-app-frontend-1 Stopped 0.4s
 - Container ev-charging-app-user-service-1 Stopping 3.3s
 - Container ev-charging-app-station-service-1 Stopping 3.3s
 - Container ev-charging-app-booking-service-1 Stopping 3.3s
 - Container ev-charging-app-auth-service-1 Stopping 3.3s

2025-06-04 13:36:17 Frontend-1 2025/06/04 08:06:17 [notice] 1#1: worker p
```

# 5 Apache Deployment using Ansible

## Steps

1. *Setup SSH connectivity (passwordless SSH) between the control node and managed node.*
2. *Install Ansible on the control node.*
3. *Create an inventory file (/etc/ansible/hosts) with managed node IP.*
4. *Create a project directory: ~/ansible-apache-nginx*
5. *Create a playbook file deploy\_apache.yml with the following content:*

```

```

```
- name: Deploy Apache2 with 4 customized HTML pages
 hosts: webservers
 become: yes
```

```
 tasks:
 - name: Install Apache2
 apt:
 name: apache2
 state: present
 update_cache: yes
```

```
- name: Ensure Apache listens on port 80
 lineinfile:
 path: /etc/apache2/ports.conf
 regexp: '^Listen '
 line: 'Listen 80'
 notify: Restart Apache

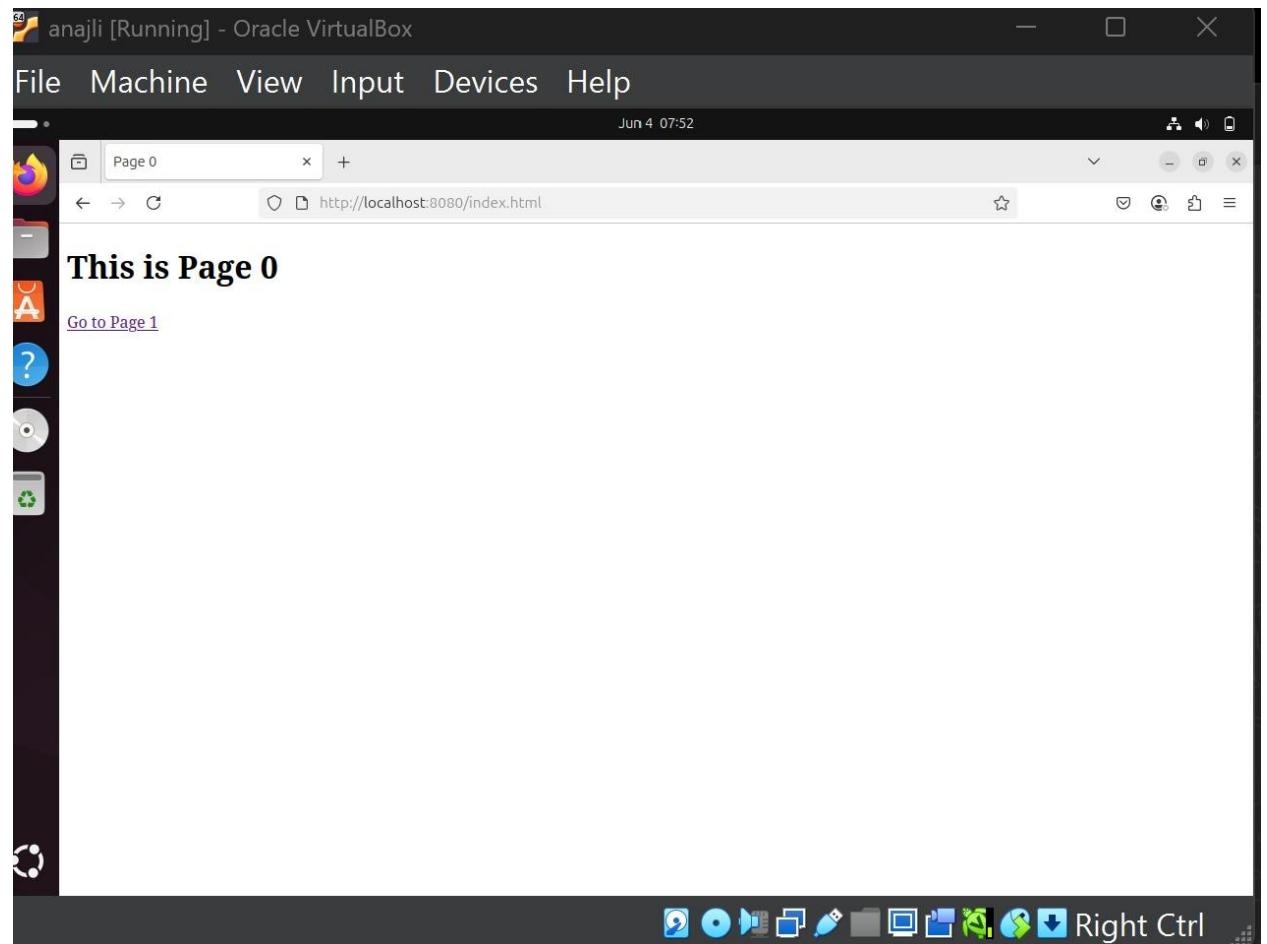
- name: Deploy 4 customized HTML pages
 copy:
 dest: "/var/www/html/page{{ item }}.html"
 content: |
 <html>
 <head><title>Apache Page {{ item }}</title></head>
 <body><h1>Apache Page {{ item }}</h1><p>Deployed by Ansible</p></body>
 </html>
 mode: '0644'
 loop: [1, 2, 3, 4]
```

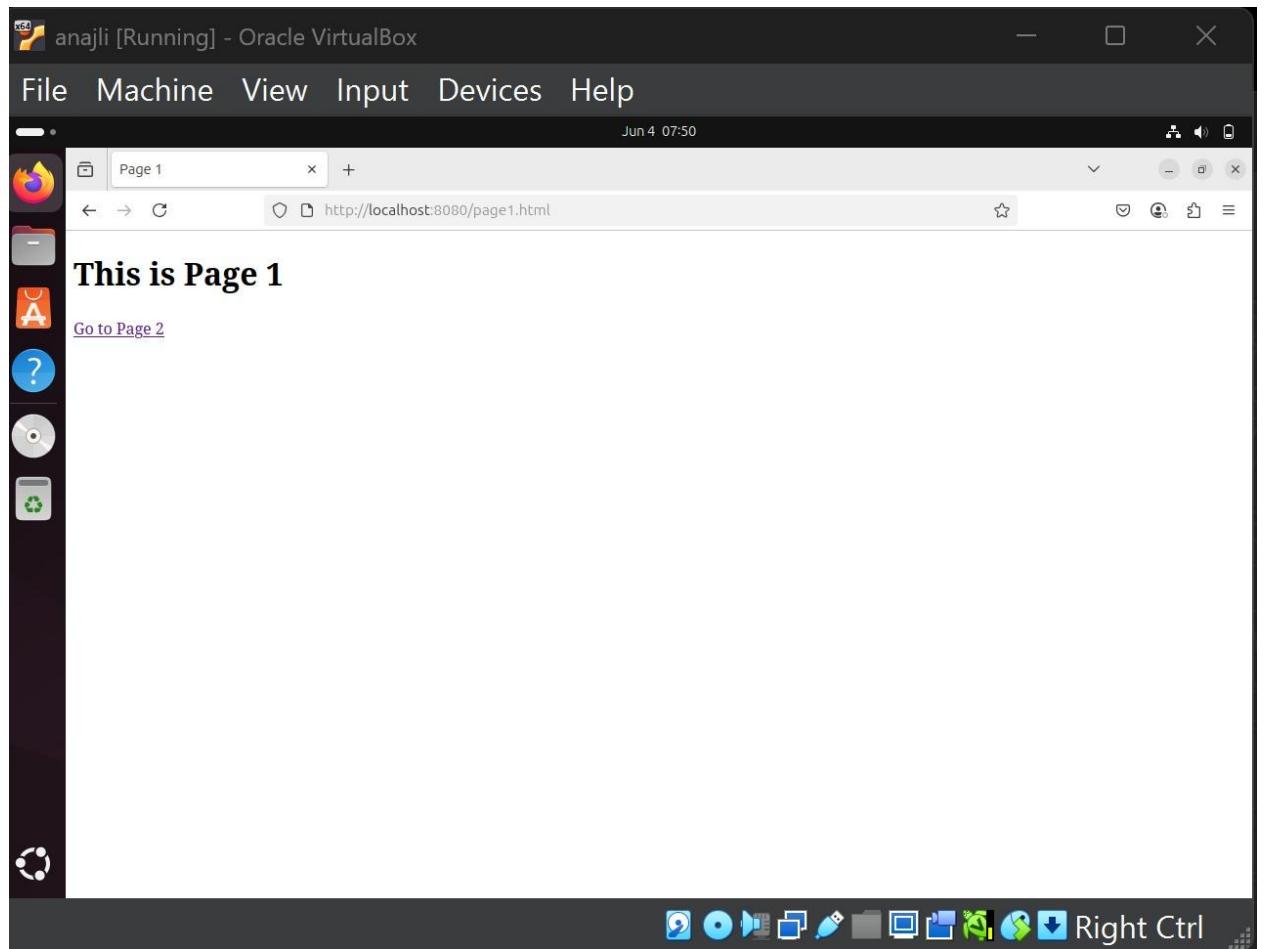
#### handlers:

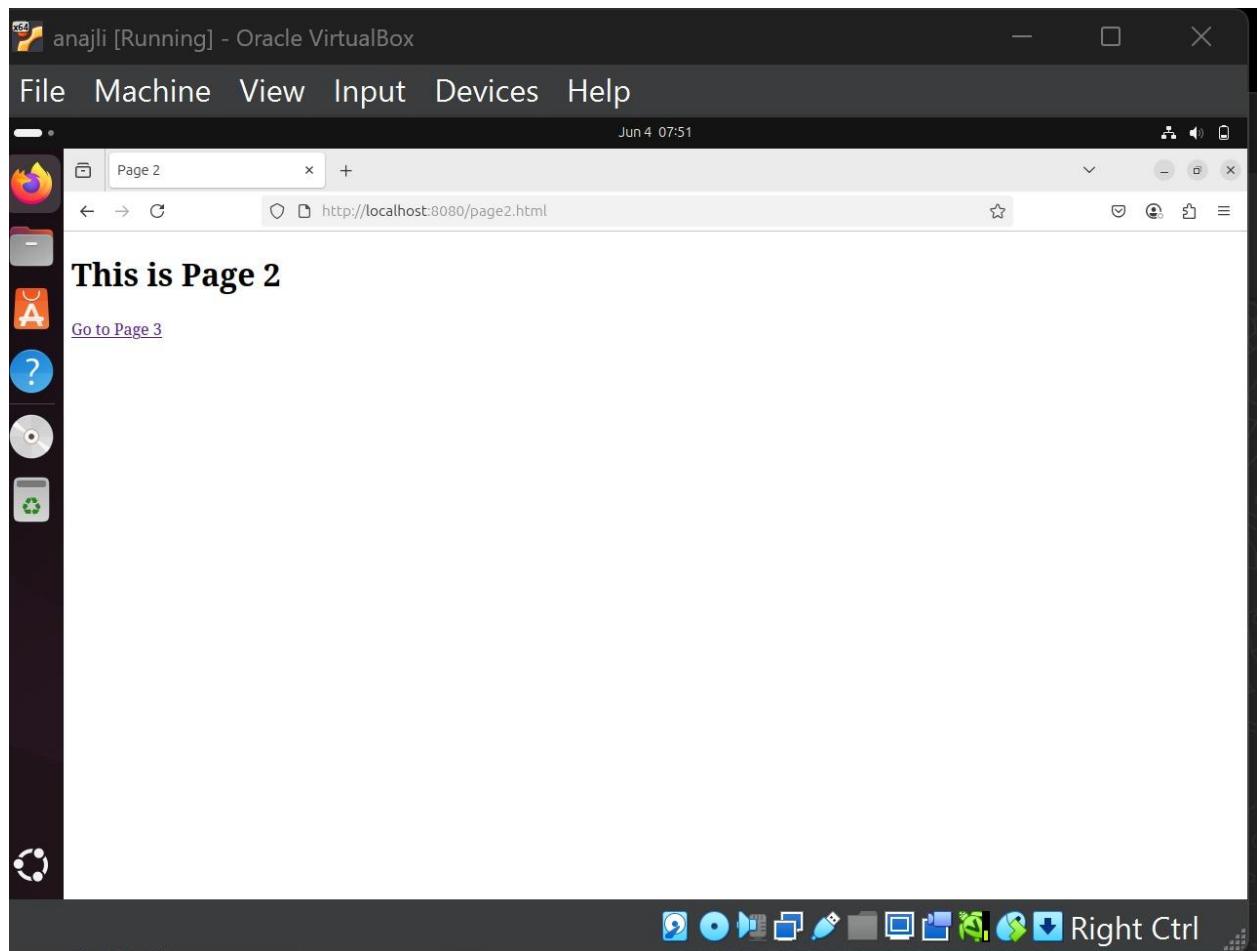
```
- name: Restart Apache
 service:
 name: apache2
 state: restarted
```

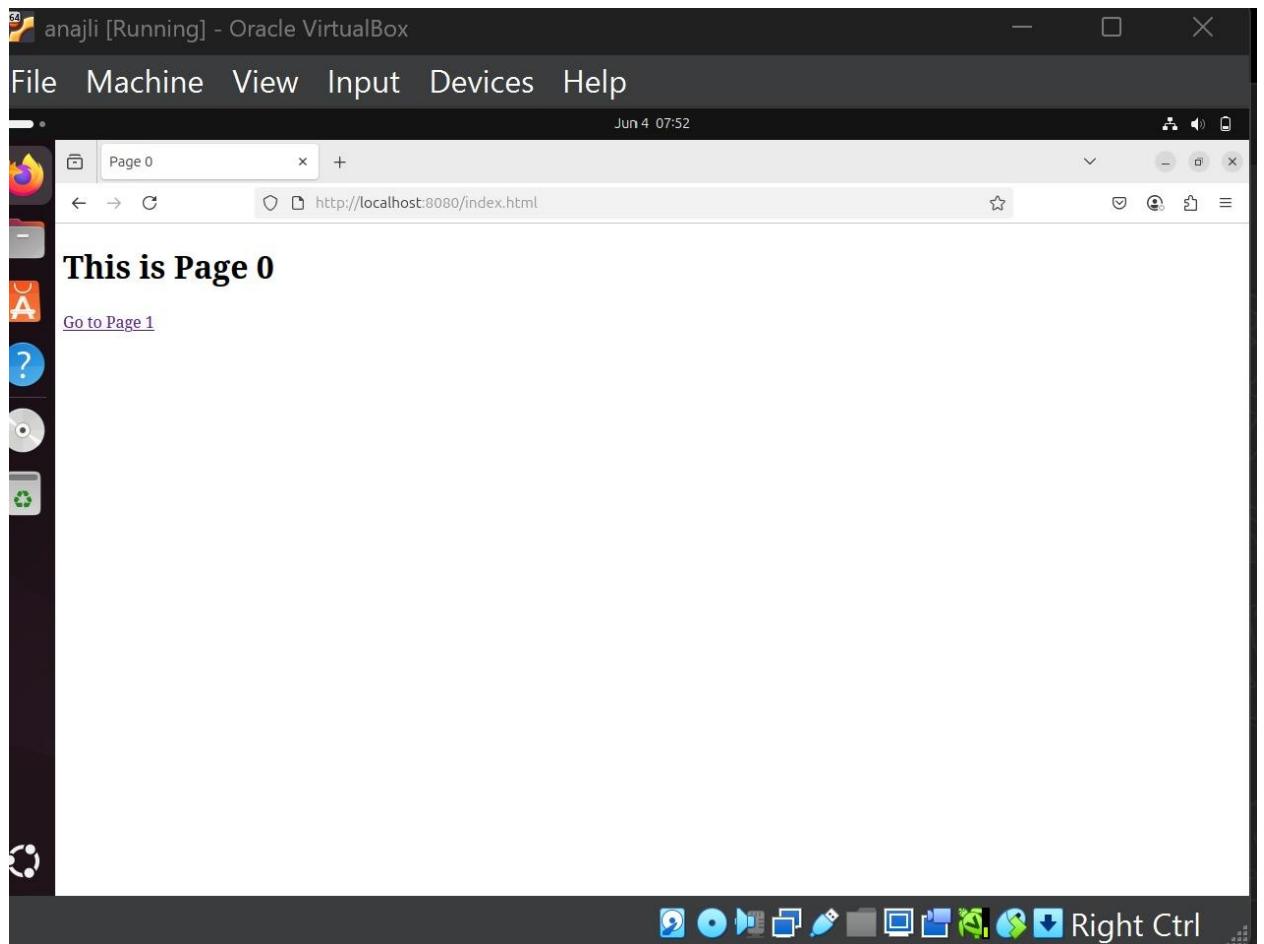
#### 5. Run the playbook:

```
ansible-playbook -i /etc/ansible/hosts deploy_apache.yml
```









7. Access pages via browser: `http://<managed-node-ip>/page1.html`

# Nginx Deployment using Ansible

## Steps

1. Setup SSH connectivity (passwordless SSH) between the control node and managed node.
2. Install Ansible on the control node.
3. Create an inventory file (/etc/ansible/hosts) with managed node IP.
4. Create a project directory: ~/ansible-apache-nginx
5. Create a playbook file deploy\_nginx.yml with the following content:

---

```

- name: Deploy Nginx with 4 customized HTML pages
 hosts: webservers
 become: yes

 tasks:
 - name: Install Nginx
 apt:
 name: nginx
 state: present

 - name: Configure Nginx to listen on port 100
 lineinfile:
 path: /etc/nginx/sites-available/default
 regexp: 'listen '
 line: ' listen 100 default_server;'
 notify: Restart Nginx

 - name: Deploy 4 customized HTML pages
 copy:
 dest: "/var/www/html/nginx_page{{ item }}.html"
 content: |
 <html>
 <head><title>Nginx Page {{ item }}</title></head>
 <body><h1>Nginx Page {{ item }}</h1><p>Deployed by Ansible</p></body>
 </html>
 mode: '0644'
 loop: [1, 2, 3, 4]

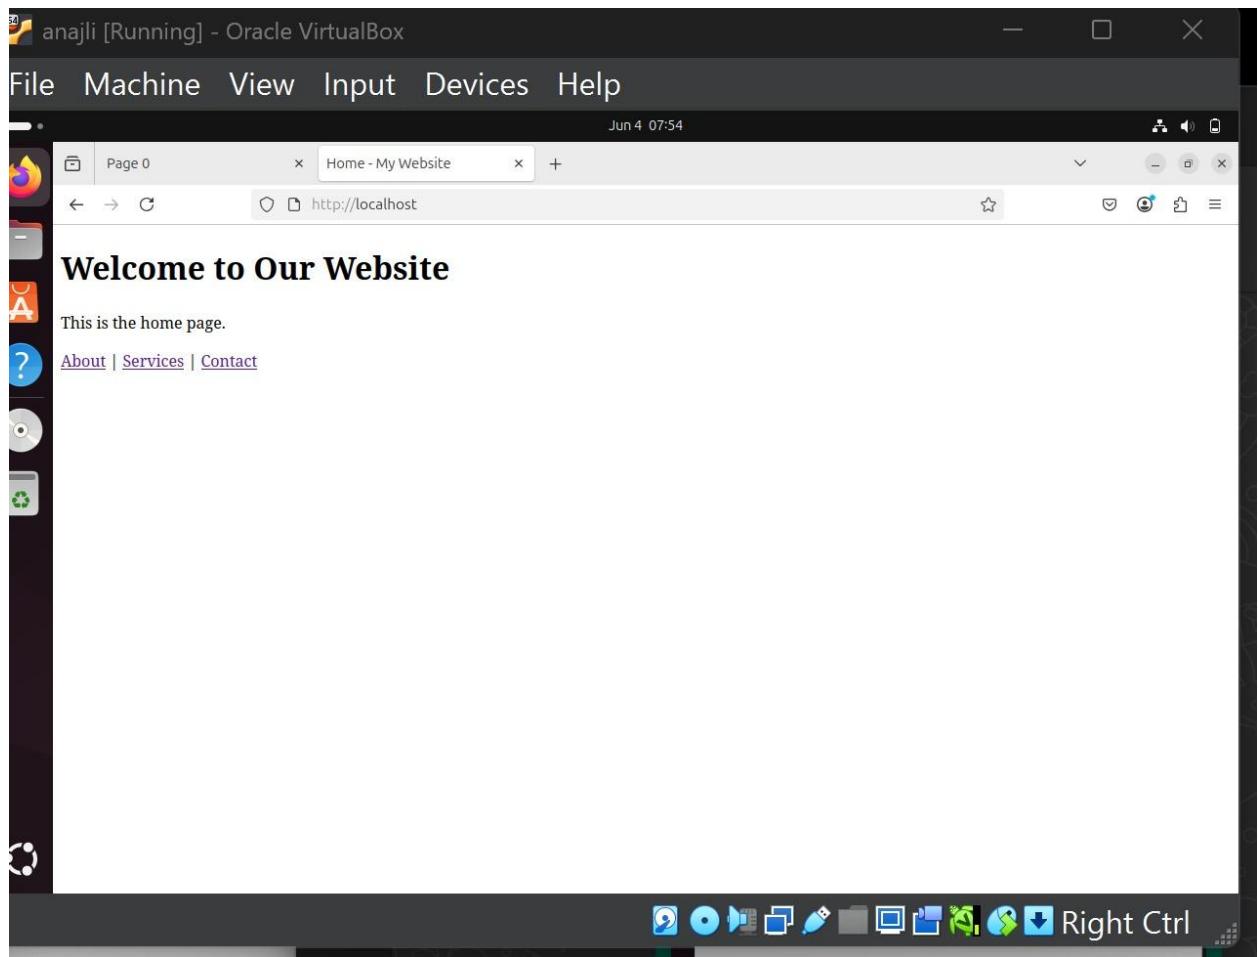
```

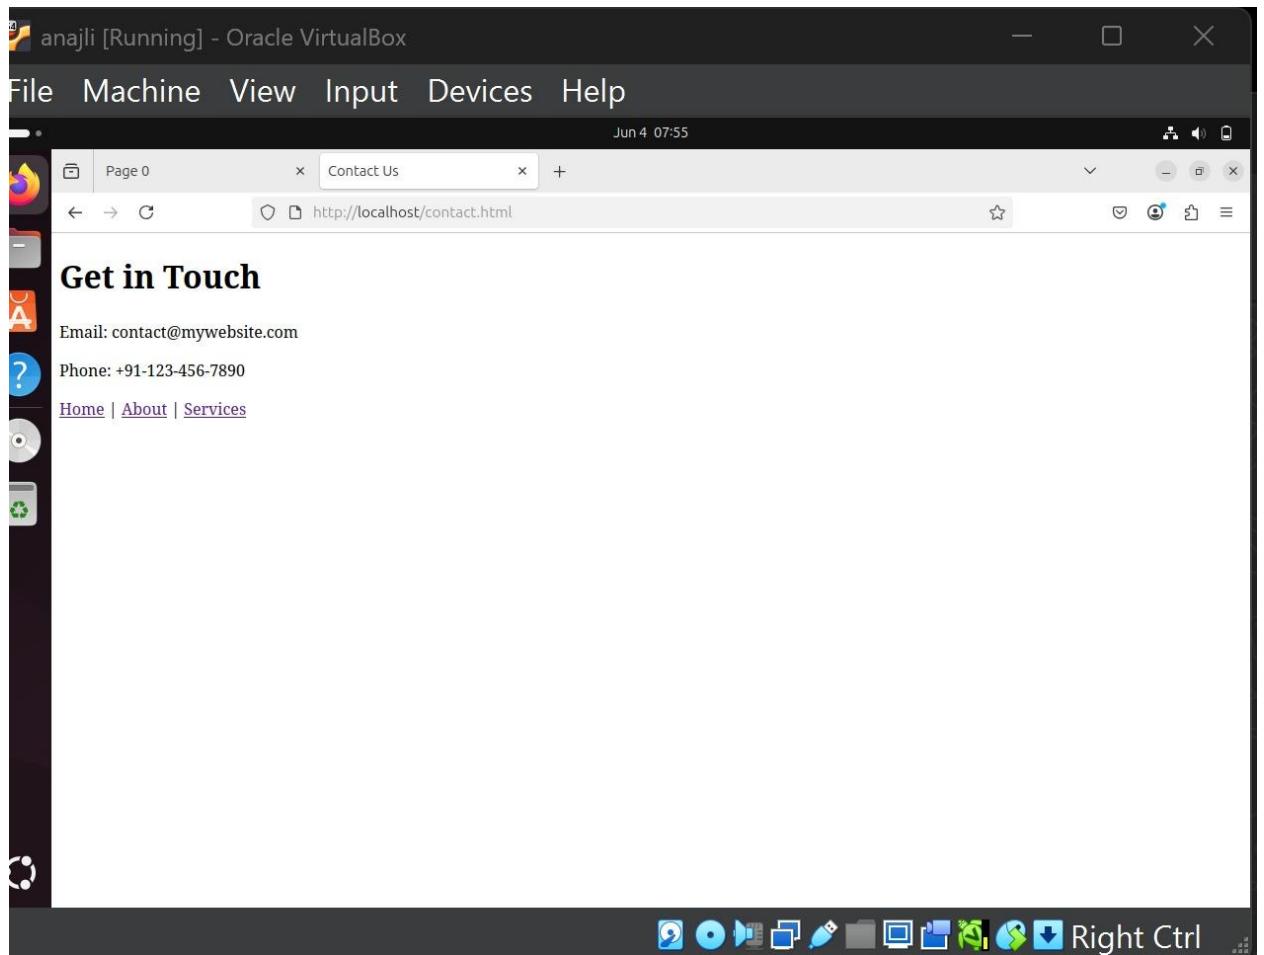
handlers:

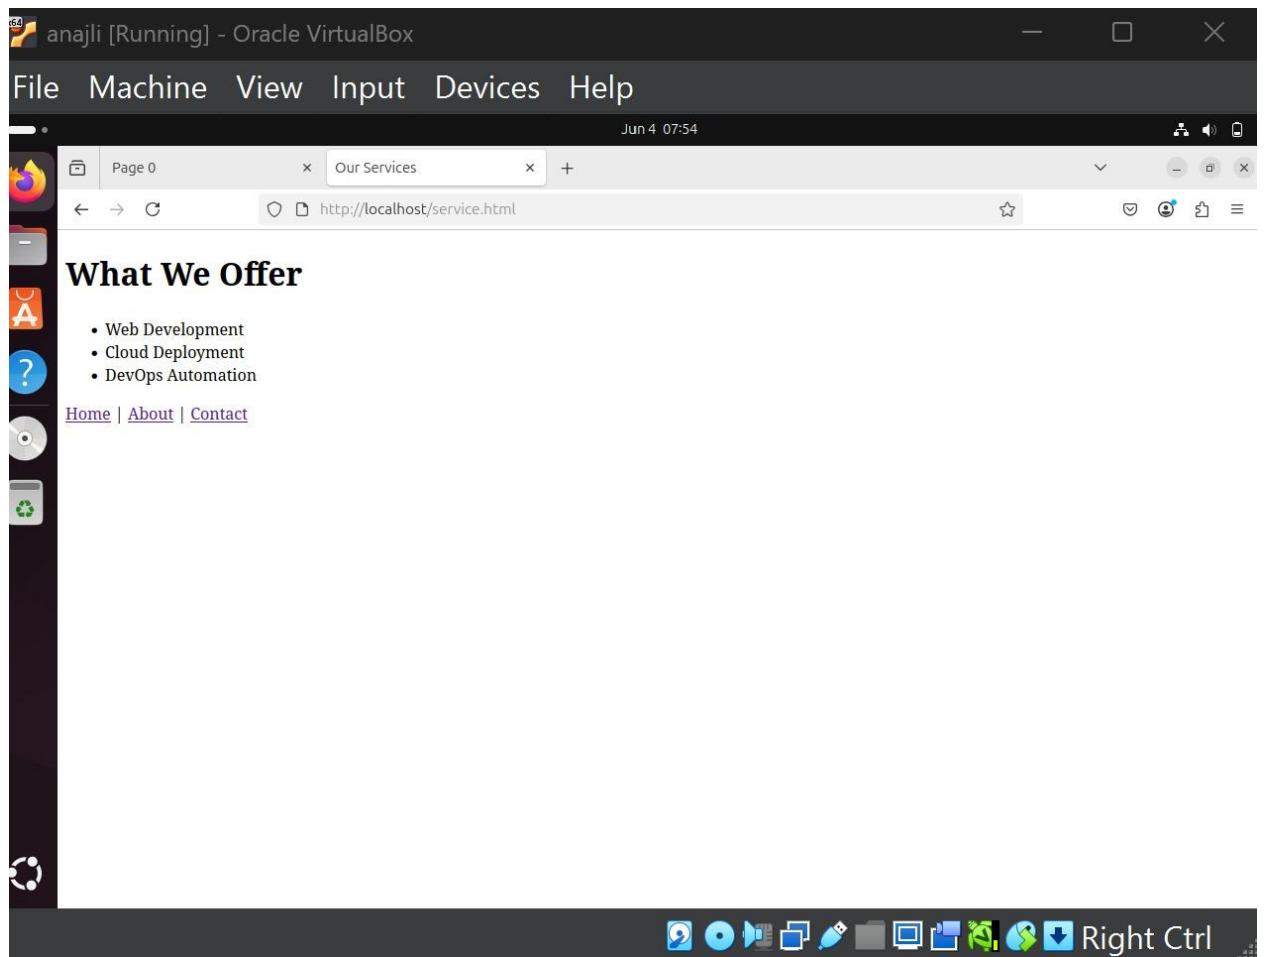
```
- name: Restart Nginx
 service:
 name: nginx
 state: restarted
```

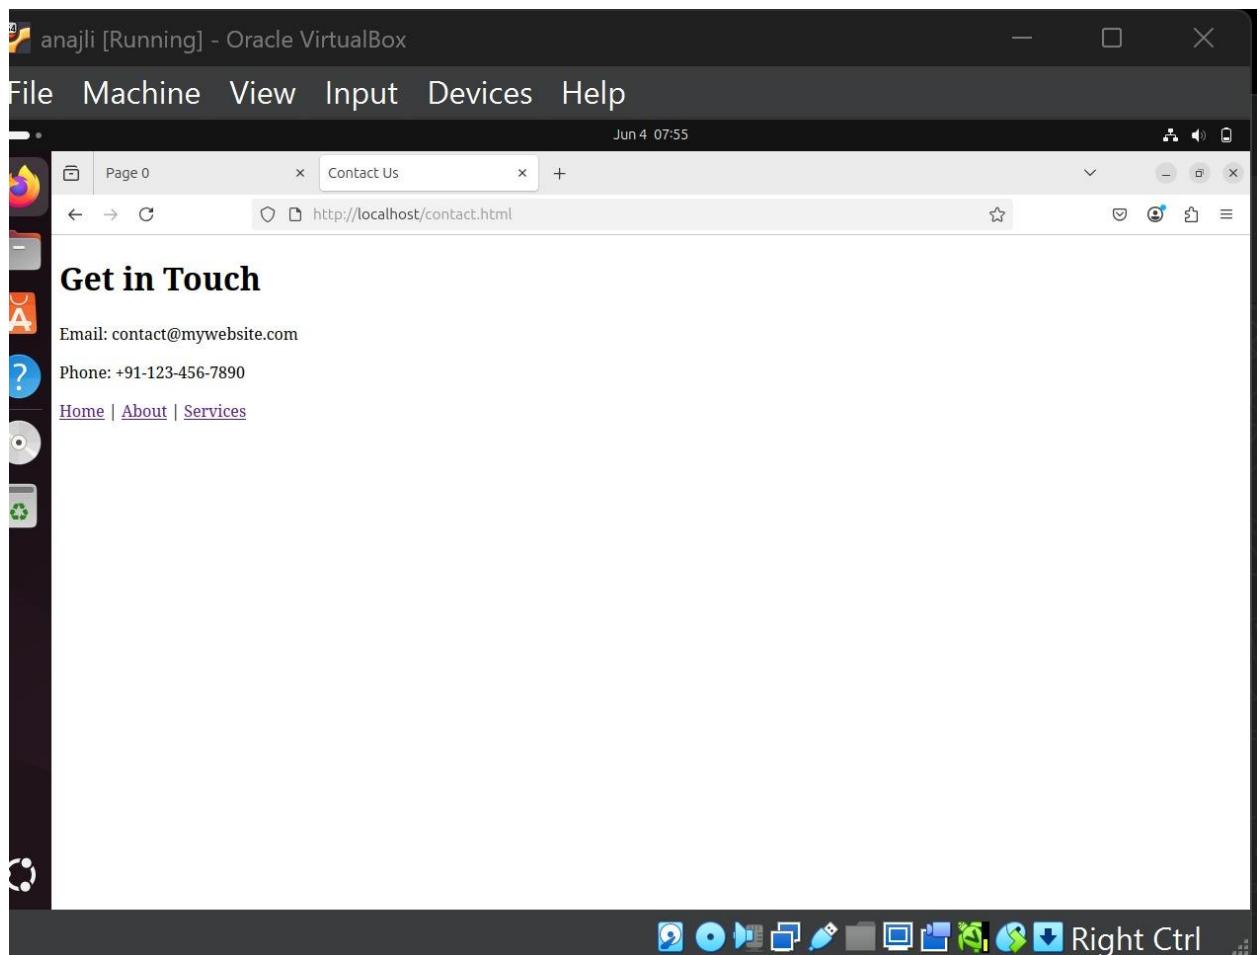
7. Run the playbook:

```
ansible-playbook -i /etc/ansible/hosts deploy_nginx.yml
```









7. Access pages via browser: [http://<managed-node-ip>:100/nginx\\_page1.html](http://<managed-node-ip>:100/nginx_page1.html)